

Dipartimento Matematica

LECTURE NOTES FOR STATISTICS

Professore
Stefano Favaro

Laureando
Claudio Meggio

Academic Year 2018/2019

Contents

| | |
|---|-----------|
| Sommario | 2 |
| 1 Introduction | 3 |
| 1.1 Inequalities | 3 |
| 1.2 Common distributions | 3 |
| 1.3 Linear Algebra | 5 |
| 1.4 Multivariate analysis | 6 |
| 1.5 Basic Concepts of Random Samples | 8 |
| 2 Concentration Measure | 11 |
| 2.1 Concentration for sum of random variables | 12 |
| 3 Likelihood Function | 17 |
| 3.1 Likelihood principles | 17 |
| 3.2 Condition of Regularity | 18 |
| 3.3 Properties of the Likelihood Function | 19 |
| 3.4 Exponential Families | 20 |
| 3.5 Natural Exponential Families | 22 |
| 4 Statistics | 23 |
| 4.1 Estimators | 27 |
| 4.2 Properties of Estimators | 28 |
| 4.3 Method of Moments | 32 |
| 4.4 Maximum Likelihood Approach | 33 |
| 5 Hypothesis Testing | 37 |
| 6 Exercises | 45 |
| Bibliografia | 51 |

Informazioni sul corso

mail: stefano.favaro@unito.it

book: [1]

1 Introduction

1.1 Inequalities

Markow's Inequality:

Let Y be a non negative random variable with finite expected value then

$$\mathbb{P}(Y \geq t) \leq \frac{\mathbb{E}[Y]}{t}$$

Chebyshev's Inequality:

Let X be a random variable with finite second moment and let $\sigma = \sqrt{\text{Var}(x)}$, then for any positive real h

$$\mathbb{P}(|X - \mathbb{E}[X]| \geq h\sigma) \leq \frac{1}{h^2} \quad (1.1.1)$$

Theorem 1.1.1. Schwarz If X, Y are random variables with finite second moment then:

$$(\mathbb{E}[XY])^2 \leq \mathbb{E}[X^2]\mathbb{E}[Y^2]$$

If X is a random variable taking values in a set I with $\mathbb{E}[X] = \mu$ and $f(\cdot)$ is convex on I , then $f(x) \geq f(\mu) + h(x - \mu)$ holds with probability 1 for some choice of h . By integrating both sides of the inequality with respect to the distribution of X we obtain the **Jensen's Inequality**:

$$\mathbb{E}[f(x)] \geq f(\mathbb{E}[X])$$

1.2 Common distributions

Gaussian:

A continuous random variable Y is said to have a Gaussian distribution with parameters μ and σ^2 if the density function at t is:

$$\frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(x - \mu)^2}{2\sigma^2} \right\}$$

Y is unimodal and symmetric around the mode $t = \mu$ and we write

$$Y \sim N(\mu, \sigma^2)$$

It's *characteristic function* is

$$\mathbb{E}[e^{tiy}] = \exp \left\{ it\mu - \frac{\sigma^2 t^2}{2} \right\}$$

The derivative of the characteristic function valued at $t = 0$ give us the non centered moments of Y .

If $Y \sim N(\mu\sigma^2)$ and $a, b \in \mathbb{R}$ then $(a + bY) \sim N(a + b\mu, b^2\sigma^2)$. This means that the entire family of distribution can be generated by linear transformations starting from any member of the family (i.e. is a *location scale family*)

If $Y_1 \sim N(\mu_1, \sigma_1^2)$ and $Y_2 \sim N(\mu_2, \sigma_2^2)$ and $Y_1 \perp Y_2$ then $Y_1 + Y_2 \sim N(\mu_1 + \mu_2, \sigma_1^2 + \sigma_2^2)$. This result can be extended to linear combinations of Gaussian random variables.

Uniform distribution:

A continuous random variable Y with density function:

$$f(t; a, b) = \frac{1}{b-a} \mathbb{1}_{[a, b]}(t)$$

is said to be *uniformly distributed* in $[a, b]$ and we write $Y \sim U(, b)$.

Theorem 1.2.1. Integral transformation theorem

If Z is a continuous random variable with distribution function F then the random variable

$$W := F(Z) \sim U(0, 1)$$

Proof.

$$\begin{aligned} \mathbb{P}(W \leq t) &= \mathbb{P}(F(Z) \leq t) \\ &= \mathbb{P}(Z \leq F^{-1}(t)) \\ &= F(F^{-1}(t)) \\ &= t \end{aligned}$$

Which is the distribution function of uniform in $[0, 1]$

□

Gamma distribution:

The *Gamma function* is:

$$\Gamma(x) := \int_0^{+\infty} t^{x-1} e^{-t} dt$$

Some properties of this function are:

- $\Gamma(x+1) = x\Gamma(x)$
- if x is a positive integer $\Gamma(x) = (x-1)!$
- $\Gamma(1) = 1$
- $\Gamma\left(-\frac{1}{2}\right) = \sqrt{\pi}$

Stirling's Approximation

$$n! \sim \sqrt{2\pi n} \left(\frac{n}{e}\right)^n$$

Definition 1.2.1. We say that a continuous random variable X has a *Gamma distribution* with shape parameter w and scale parameter λ ($X \sim \text{Gamma}(w, \lambda)$), if its density function is:

$$f(t; w, \lambda) = \frac{\lambda^w}{\Gamma(w)} t^{w-1} e^{-\lambda t} \mathbb{1}_{[\mathbb{R}^+]}(t)$$

Proposition 1.2.1. If $Y_1 \sim \text{Gamma}(w_1, \lambda)$ and $Y_2 \sim \text{Gamma}(w_2, \lambda)$ and $Y_1 \perp\!\!\!\perp Y_2$ then:

$$Y_1 + Y_2 \sim \text{Gamma}(w_1 + w_2, \lambda)$$

Other distributions:

1. Beta distribution
2. Binomial Distribution
3. Hypergeometric Distribution
4. Negative Binomial distribution

1.3 Linear Algebra

Matrix:

Consider A, B two $n \times n$ squared matrix

Notation:

I_n is the identity matrix of order n

1_n is the $n \times 1$ (column) vector with all elements equal to 1

\bigcirc is a matrix with all element equal to zero

$|A|$ denotes the determinant of A

Definitions / Properties:

Definition 1.3.1. A is called *symmetric matrix* if

$$A = A^T$$

Proposition 1.3.1. for two conformable matrix A, B we have:

$$|AB| = |A||B|$$

Definition 1.3.2. if $|A| \neq 0$, A is called *non singular* or *invertible* and there exist a matrix A^{-1} called *inverse* such that

$$AA^{-1} = A^{-1}A = I_n$$

Definition 1.3.3. A *diagonal matrix* is a matrix with all elements outside the main diagonal equal to zero

Definition 1.3.4. A matrix A is called *invertible* if there exist an invertible matrix P such that $P^{-1}AP$ is a diagonal matrix

Proposition 1.3.2. It holds:

$$(A^T)^{-1} = (A^{-1})^T$$

$$(AB)^{-1} = B^{-1}A^{-1}$$

Definition 1.3.5. A symmetric matrix A is said to be *positive semi-definite* if

$$v^T Av \geq 0, \forall v \in \mathbb{R}^n$$

Definition 1.3.6. A matrix A is called *orthogonal* if:

$$A^{-1} = A^T$$

Definition 1.3.7. Given a matrix A , we call the sum of all the elements on the main diagonal *trace* of A :

$$Tr(A) := \sum_{i=1}^n a_{ii}$$

Proposition 1.3.3. For any matrix A, B we have:

$$Tr(AB) = Tr(BA)$$

Definition 1.3.8. An *idempotent matrix* A is a matrix which, when multiplied by itself, yields itself i.e.:

$$AA = A$$

Proposition 1.3.4. Properties of an idempotent matrix A :

1. $I - A$ is also an idempotent matrix
2. A is idempotent if and only if for all positive integers k , $A^k = A$
3. an idempotent matrix is always diagonalizable and its eigenvalues are either 0 or 1
4. the trace of an idempotent matrix is always an integer and equal to its rank

Proposition 1.3.5. Two identities:

1. $(A + BCD)^{-1} = A^{-1} - A^{-1}B(C^{-1} + DA^{-1}B)^{-1}DA^{-1}$ **Binomial Jensen Theorem**
2. $(A + bd^T)^{-1} = A^{-1} - \frac{1}{1+d^T A^{-1}b} A^{-1}bd^T A^{-1}$

Theorem 1.3.1. Spectral theorem

Let A be a symmetric $n \times n$ matrix, then there exists an orthogonal matrix Q such that:

$$A = Q\Lambda Q^T$$

where Λ is a diagonal matrix whose diagonal elements are the eigenvalues $\lambda_1, \dots, \lambda_n$ of A .

Corollary 1.3.1.

$$|A| = |\Lambda| = \prod_{i=1}^n \lambda_i$$

1.4 Multivariate analysis

Definition 1.4.1. Take X_1, \dots, X_n random variables defined on the same probability space, we define the *random vector* or the *multivariate random variable* X as:

$$X = (x_1, \dots, x_n)^T$$

Definition 1.4.2. The *mean vector* of X is obtained by forming the vector of the mean values of the components

$$\mathbb{E}[X] = (\mathbb{E}[X_1] \dots \mathbb{E}[x_n])^T$$

Definition 1.4.3. Similarly we can define the *variance matrix* as:

$$Var[X] := \begin{bmatrix} Var(X_1) & Cov(X_1, X_2) & \dots & Cov(X_1, X_n) \\ Cov(X_2, X_1) & Var(X_2) & \dots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ Cov(X_n, X_1) & \dots & \dots & Var(X_n) \end{bmatrix}$$

Definition 1.4.4. A generic element of the *correlation matrix* is defined as following:

$$Corr(x_i, x_j) := \frac{Cov(X_i, X_j)}{\sqrt{Var(x_i)Var(X_j)}}$$

Lemma 1.4.1. Let $A = a_{ij}$ be a $k \times n$ matrix, $b = (b_1 \dots b_n)^T$ a $n \times 1$ vector and $x = (X_1 \dots X_n)$ a random vector with $\mathbb{E}[x] = \mu$, $Var(x) = V$ define

$$Y := Ax + b$$

then

$$\mathbb{E}[Y] = A\mu + b$$

$$Var[Y] = AVA^T$$

Lemma 1.4.2. The variance matrix V of the random vector X is positive semi-definite if there exists no vector b , such that b^T is a degenerate random variable.

Lemma 1.4.3. If $A = (a_{ij})$ is a $n \times n$ matrix then:

$$\mathbb{E}[X^T A X] = \mu^T A \mu + \text{Tr}(AV)$$

where $V = \text{Var}(X)$

Proof.

$$\begin{aligned} \mathbb{E}[X^T A X] &= \mathbb{E}\left[\sum_{i=1}^n \sum_{j=1}^n x_i a_{ij} x_j\right] \\ &= \sum_{i=1}^n \sum_{j=1}^n a_{ij} \mathbb{E}[x_i x_j] \\ &= \sum_{i=1}^n \sum_{j=1}^n a_{ij} \mu_i \mu_j + v_{ij} \\ &= \mu^T A \mu + \sum_{i=1}^n (AV)_{ii} \\ &= \mu^T A \mu + \text{Tr}(AV) \end{aligned}$$

□

Multivariate Gaussian distribution:

Consider a vector $Z = (Z_1 \dots Z_k)^T$ where $Z_1 \dots Z_k$ are independent and identically distributed standard Gaussian random variables. Now set

$$Y = AZ + \mu$$

Where A is a non singular $k \times k$ matrix and μ is a $k \times 1$ vector.

It is natural to define Y as a *k-generated distribution of the Gaussian distribution*.

We start from:

$$f_z = \frac{1}{(2\pi)^{k/2}} \exp \left\{ -\frac{1}{2} t^T t \right\}$$

(Z are independent so we simply multiplied them).

Since $Z = A^{-1}(Y - \mu)$, the Jacobian of the transformation is:

$$\left| \frac{dz_i}{dy_i} \right| = |A|^{-1} = |V|^{-1/2}$$

Taking into account that $|V| = |AA^T| = |A|^2$.

Setting $Y = At + \mu \implies t = A^{-1}(Y - \mu)$ we obtain:

$$t^T t = \{A^{-1}(Y - \mu)\}^T \{A^{-1}(Y - \mu)\} = (Y - \mu)^T V^{-1} (Y - \mu)$$

Therefore the density of Y is:

$$f_Y(y) = \frac{1}{(2\pi)^{k/2} |V|^{1/2}} \exp \left\{ -\frac{1}{2} (y - \mu)^T V^{-1} (y - \mu) \right\}$$

We say that the random variable $Y = (Y_1 \dots Y_n)^T$ with density function f_Y is a multivariate Gaussian random variable with mean μ and variance V . $Y \sim N_k(\mu, V)$.

Now we will explore *marginal and conditional* distribution of Y .

Proposition 1.4.1. *If A is a $k \times k$ positive matrix and b is a $k \times 1$ vector then:*

$$\int_{\mathbb{R}^k} \frac{1}{(2\pi)^{k/2}} \exp \left\{ -\frac{1}{2}(y^T A y - 2b^T y) \right\} dy = \frac{\exp\{1/2 b^T A^{-1} b\}}{|A|^{1/2}}$$

Proof. Let $\mu A^{-1}b$ and within the integral expand \exp by adding and subtracting $\frac{1}{2}\mu^T A \mu$ so that

$$\int_{\mathbb{R}^k} \frac{1}{(2\pi)^{k/2}} \exp \left\{ -\frac{1}{2}(y^T A y - 2b^T y) \right\} dy = |A^{-1}|^{1/2} \exp \left\{ \frac{1}{2}\mu^T A \mu \right\} \int_{\mathbb{R}^k} g(y)$$

□

1.5 Basic Concepts of Random Samples

Definition 1.5.1. Let $X_1 \dots X_n$ independent and identically distributed random variables with distribution $\sim f_{X_i}(x_i; \theta)$. We call $X := (X_1 \dots X_n)$ *random sample*.

According with the definition of random sample the distribution of X will be:

$$f_X(X; \theta) = \prod_{i=1}^n f_{x_i}(x_i; \theta)$$

Definition 1.5.2. We denote by $x = (x_1 \dots x_n)$ the observed sample

Definition 1.5.3. A *statistical model* is defined as following

$$\{f_X(x; \theta) : \theta \in \Theta\}$$

Where Θ is the *parametric space*

Usually Θ will be a open subset of \mathbb{R}^n .

Definition 1.5.4. Let X_1, \dots, X_n be a random sample of size n from a population and let $T(x_1, \dots, x_n)$ be a real-valued or vector-valued function whose domain includes the sample space of (X_1, \dots, X_n) . Then the random variable or random vector $T_n = T(X_1, \dots, X_n)$ is called a *statistic*. The probability distribution of a statistic T_n is called the *sampling distribution* of T_n .

Some examples of statistic are:

- Sample mean: $\bar{X}_n := \frac{1}{n} \sum_{i=1}^n x_i$
- Sample variance: $\tilde{S}^2 := \frac{1}{n} \sum_{i=1}^n (x_i - \bar{X})^2$
- Corrected sample variance: $S^2 := \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{X})^2$
- Sample moments of order r : $M_{r,n} := \frac{1}{n} \sum_{i=1}^n x_i^r$
- Sample moments of order r : $\bar{M}_{r,n} := \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x}_n)^r$
- Ordered statistic: $X_{(m)}$
- Sample min: $X_{(1)}$
- Sample max: $X_{(n)}$
- Sample median: $Me := \begin{cases} \frac{1}{2}(X_{(n/2)} + X_{(n/2+1)}) & \text{if } n - \text{even} \\ X_{(\frac{n+1}{2})} & \text{if } n - \text{odd} \end{cases}$

Finding the distribution of T_n in general it is complex. We can make it easier by giving constraints. Suppose for example $X \sim N(\mu, \sigma^2) \leftarrow$ fair assumption because there is the Central Limit Theorem.

Theorem 1.5.1. Fisher-Cochran

Let Q, Q_1, Q_2 random variables such that $Q = Q_1 + Q_2$ and let $Q \sim \chi_g^2$ and $Q_2 \sim \chi_{g_1}^2$. Then

$$Q_2 \sim \chi_{g_2}^2 \quad \text{where } g_2 = g - g_1,$$

and $Q_1 \perp\!\!\!\perp Q_2$

Proposition 1.5.1. Let $X_i \sim N(\mu, \sigma^2)$ and $X = (X_1 \dots X_n)$. Then

1. $\bar{X}_n := \frac{1}{n} \sum_{i=1}^n x_i \sim N(\mu, \frac{\sigma^2}{n})$
2. $\tilde{S}_n := \frac{1}{n} \sum_{i=1}^n (x_i - \bar{X})^2 \sim \frac{\sigma^2}{n} \chi_{n-1}^2$

where χ_{n-1}^2 is the Chi-squared distribution with $n - 1$ degrees of freedom.

Proof. 1. the first one is easily checked using the linearity of the Gaussian distribution.

2. Consider the random variable $\tilde{S}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \mu)^2$ and proceed as following:

$$\begin{aligned} \frac{n\tilde{S}^2}{\sigma^2} &= \sum_{i=1}^n \left(\frac{x_i - \mu}{\sigma} \right)^2 \\ &= \sum_{i=1}^n \left(\frac{x_i - \mu + \bar{x}_n - \bar{x}_n}{\sigma} \right)^2 \\ &= \sum_{i=1}^n \left(\frac{x_i - \bar{x}_n}{\sigma} \right)^2 + \sum_{i=1}^n \left(\frac{\bar{x}_n - \mu}{\sigma} \right)^2 + 2 \sum_{i=1}^n \left(\frac{x_i - \bar{x}_n}{\sigma} \right) \left(\frac{\bar{x}_n - \mu}{\sigma} \right) \end{aligned}$$

Now consider separately the three terms of the sum:

$$\sum_{i=1}^n \left(\frac{x_i - \bar{x}_n}{\sigma} \right) = n\bar{x}_n \sum_{i=1}^n x_i = n\bar{x}_n - n\bar{x}_n = 0$$

$$\implies 2 \sum_{i=1}^n \left(\frac{x_i - \bar{x}_n}{\sigma} \right) \left(\frac{\bar{x}_n - \mu}{\sigma} \right) = 0$$

For the other two terms consider:

$$\sum_{i=1}^n \left(\frac{x_i - \mu}{\sigma} \right)^2 \sim \chi_1^2$$

$$\text{And } \sum_{i=1}^n \left(\frac{\bar{x}_n - \mu}{\sigma} \right)^2 = \left(\frac{\bar{x}_n - \mu}{\sigma/\sqrt{n}} \right)^2 \sim \chi_n^2$$

So, using the theorem 1.5.1

$$\sum_{i=1}^n \left(\frac{x_i - \mu}{\sigma} \right)^2 = \sum_{i=1}^n \left(\frac{x_i - \bar{x}_n}{\sigma} \right)^2 + \left(\frac{\bar{x}_n - \mu}{\sigma/\sqrt{n}} \right)^2 \sim \chi_{n-1}^2$$

$$\implies \tilde{S}_n^2 \sim \frac{\sigma^2}{n} \chi_{n-1}^2$$

□

Let $X = (X_1 \dots X_n)$ be a random sample where $X_i \sim N(\mu, \sigma^2)$, consider the statistic:

$$T_n = \frac{\bar{X}_n - \mu}{S_n/\sqrt{n}} = \frac{(\bar{X}_n - \mu)/\sigma/\sqrt{n}}{\sqrt{S_n/\sigma^2}} = \frac{Z}{\sqrt{R/(n-1)}}$$

where $Z \sim N(0, 1)$, $R \sim (\chi_n - 1)$. We can say that $T_n \sim T$ - Student with $(n - 1)$ degrees of freedom only if $Z \perp\!\!\!\perp R$. We can they are independent because of the following:

Theorem 1.5.2. If $(x_1 \dots x_n)$ is a random sample with $X_i \sim N(\mu, \sigma^2)$, then

$$\bar{X}_n = \frac{1}{n} \sum_{i=1}^n x_i$$

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{X}_n)^2$$

are independent

The other way around is also true:

Theorem 1.5.3. If $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n x_i$, $S^2 = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{X}_n)^2$ are independent, then $X = (x_1 \dots x_n)$ is random sample, where $X_i \sim N(\mu, \sigma^2)$.

2 Concentration Measure

We're now going to investigate some methods to study the tail of a distribution.

Consider a non negative random variable and let $t > 0$. Then by Markow inequality we have:

$$\mathbb{P}(X \geq t) \leq \frac{\mathbb{E}[X]}{t}$$

We can try to improve this inequality using a function Φ that is strictly increasing with non negative values. Then we can write

$$\mathbb{P}(X \geq t) = \mathbb{P}(\Phi(X) \geq \Phi(t)) \leq \frac{\mathbb{E}[\Phi(X)]}{\Phi(t)}$$

In particular we can take $\Phi(x) = x^q$, $X \geq 0, q > 0$ so we have

$$\mathbb{P}(|X - \mathbb{E}[X]| \geq t) \leq \frac{\mathbb{E}[|X - \mathbb{E}[X]|^q]}{t^q}$$

In specific examples one can choose the value of q that optimises the upper bound.

A related idea is at the basis of **Chernoff's bounding method**: taking $\Phi(X) = e^{sX}$ where s is an arbitrary positive number for any random variable X and $t \in \mathbb{R}$ we have:

$$\mathbb{P}(X \geq t) = \mathbb{P}(e^{sX} \geq e^{st}) \leq \frac{\mathbb{E}[e^{sX}]}{e^{st}} \quad (2.0.1)$$

So we can bound the probability using the characteristic function which is usually easier to compute than $\mathbb{E}[X^q]$.

However it can be proven that the bounding given form $\Phi(X) = x^q$ is always better than the one given by $\Phi(X) = E^{sX}$.

Theorem 2.0.1. *Cauchy Swartz inequality*

Given two random variables with finite second moments then:

$$|\mathbb{E}[XY]|^2 \leq \mathbb{E}[X^2]\mathbb{E}[Y^2]$$

Theorem 2.0.2. *Let $t \geq 0$ then*

$$\mathbb{P}(X - \mathbb{E}[X] \geq t) \leq \frac{\text{Var}(X)}{\text{Var}(X) + t^2}$$

Proof. We assume that $\mathbb{E}[X] = 0$ (the proof for the general case is the same).

For all t we can write

$$t = \mathbb{E}[t] = \mathbb{E}[t] - \mathbb{E}[X] = \mathbb{E}[t - X] \leq \mathbb{E}[(t - X)\mathbb{1}_{[X < t]}(X)]$$

Then for $t \geq 0$ from Cauchy-Schwarz inequality:

$$\begin{aligned} t^2 &\leq \mathbb{E}[(t - X)^2]\mathbb{E}[(\mathbb{1}_{[X < t]}(X))^2] \\ &= \mathbb{E}[(t - X)^2]\mathbb{P}(X < t) \\ &= (\text{Var}(X) + t^2)\mathbb{P}(X < t) \end{aligned}$$

$$\begin{aligned} \implies \mathbb{P}(X < t) &\geq \frac{t^2}{\text{Var}(X) + t^2} \\ \implies \mathbb{P}(X \geq t) &= 1 - \mathbb{P}(X < t) \leq 1 - \frac{t^2}{\text{Var}(X) + t^2} = \frac{\text{Var}(X)}{\text{Var}(X) + t^2} \end{aligned} \quad \square$$

Theorem 2.0.3. *Let f, g be non decreasing real valued functions defined on the real line. If X is a real valued random variable then:*

$$\mathbb{E}[f(x)g(x)] \geq \mathbb{E}[f(x)]\mathbb{E}[g(x)]$$

If f is non increasing and g is non decreasing then:

$$\mathbb{E}[f(x)g(x)] \leq \mathbb{E}[f(x)]\mathbb{E}[g(x)]$$

Proof. WRONG Let Y be a random variable with the same distribution as X and $X \perp\!\!\!\perp Y$. Because f, g are non decreasing functions we have $(f(x) - f(y))(g(x) - g(y)) \geq 0$.

$$\implies 0 \leq \mathbb{E}[(f(x) - f(y))(g(x) - g(y))] = \mathbb{E}[f(x)g(x) - f(x)g(y) - f(y)g(x) + f(y)g(y)]$$

\implies

$$\begin{aligned} \mathbb{E}[f(x)g(x)] &\geq \mathbb{E}[f(x)g(y)] + \mathbb{E}[f(y)g(x)] - \mathbb{E}[f(y)g(y)] \\ &= \mathbb{E}[f(x)g(y)] \\ &= \mathbb{E}[f(x)]\mathbb{E}[g(y)] \\ &= \mathbb{E}[f(x)]\mathbb{E}[g(x)] \end{aligned}$$

The second part of the theorem can be proved in the same way. \square

The previous theorem can be generalized as following:

Theorem 2.0.4. *Let $f, g : \mathbb{R}^n \rightarrow \mathbb{R}$ be non increasing functions. Let $X_1 \dots X_n$ be independent real valued random variables and define the random variable $X = (X_1 \dots X_n)$ that take values in \mathbb{R}^n then:*

$$\mathbb{E}[f(x)g(x)] \geq \mathbb{E}[f(x)]\mathbb{E}[g(x)]$$

If f is non increasing and g is non decreasing then:

$$\mathbb{E}[f(x)g(x)] \leq \mathbb{E}[f(x)]\mathbb{E}[g(x)]$$

2.1 Concentration for sum of random variables

We want to bound the probability $\mathbb{P}(S_n - \mathbb{E}[S_n] \geq t)$ where $S_n = \sum_{i=1}^n X_i$ and $X_1 \dots X_n$ are independent random variables real valued.

An application of the Chebyshev's inequality give us:

$$\mathbb{P}(|S_n - \mathbb{E}[S_n]| \geq t) \leq \frac{\text{Var}(S_n)}{t^2} = \frac{\sum_{i=1}^n \text{Var}(X_i)}{t^2}$$

Applying the Chebyshev's inequality to $\frac{1}{n} \sum_{i=1}^n x_i$ we get

$$\begin{aligned} \mathbb{P}\left(\left|\frac{1}{n} \left(\sum_{i=1}^n x_i - \mathbb{E}[X_i]\right)\right| \geq \epsilon\right) &= \mathbb{P}\left(|S_n - \mathbb{E}[S_n]| \geq \epsilon n\right) \\ &\leq \frac{\sum_{i=1}^n \text{Var}(X_i)}{\epsilon^2 n^2} \end{aligned}$$

If we define $\sigma^2 := \frac{1}{n} \sum_{i=1}^n x_i^2$ then:

$$\mathbb{P}\left(\left|\frac{1}{n} \sum_{i=1}^n x_i - \mathbb{E}[X_i]\right| \geq \epsilon\right) \leq \frac{\sigma^2}{n\epsilon^2} \quad (2.1.1)$$

Observation 1. To understand why the equation 2.1.1 is unsatisfying recall what appens with the *Central Limit Theorem*:

$$\mathbb{P}\left(\sqrt{\frac{n}{\sigma^2}}\left(\frac{1}{n}\sum_{i=1}^n X_i - \mathbb{E}[X_i]\right) \geq y\right) \xrightarrow{n \rightarrow \infty} 1 - \Phi(y) \leq \frac{1}{\sqrt{2\pi}} \frac{e^{-y^2/2}}{y}$$

(where Φ is the CDF of the standard Gaussian distribution)

so

$$\mathbb{P}\left(\sqrt{\frac{n}{\sigma^2}}\left(\frac{1}{n}\sum_{i=1}^n X_i - \mathbb{E}[X_i]\right) \geq \epsilon\right) \lesssim \exp\left\{\frac{-n\epsilon^2}{2\sigma}\right\}$$

So for $\mathbb{P}\left(\sqrt{\frac{n}{\sigma^2}}\left(\frac{1}{n}\sum_{i=1}^n X_i - \mathbb{E}[X_i]\right) \geq \epsilon\right)$ we have:

$$\exp\left\{\frac{-n\epsilon^2}{2\sigma}\right\} \leftarrow \text{from Central Limit Theorem}$$

$$\frac{\sigma^2}{n\epsilon^2} \leftarrow \text{from Chebyshev's inequality}$$

From here we can see that the Chebyshev's inequality doesn't work well for the sum of n random variables when n is large. Meanwhile the Chebyshev's inequality works better than the Central Limit Theorem for small n .

Another instrument previously introduced that can be helpful for bounding tail probabilities of sum of independent random variables is the **Chernoff bounding** 2.0.1:

$$\mathbb{P}(S_n - \mathbb{E}[S_n] \geq t) \leq e^{-st} \mathbb{E}[\exp\{s \sum_{i=1}^n (x_i - \mathbb{E}[X_i])\}] = e^{-st} \prod_{i=1}^n \mathbb{E}[\exp\{s(x_i - \mathbb{E}[X_i])\}] \quad (2.1.2)$$

(remember that s is an arbitrary positive number)

Now the problem of finding bound on the tail probability reduces to the problem of finding (upper) bounds for the moments generating function of $X_i - \mathbb{E}[X_i]$.

As we saw Chebyshev's inequality 1.1.1 does not work well for sums of random variables. In this section we will see a partial solution given by *Hoeffding's Inequality*, then a more complete solution given by *Bernstein Inequality*.

Lemma 2.1.1. Let X be a random variable with $\mathbb{E}[X] = 0$ (actually it can be generalized for a random variable with any expected value), $a \leq X \leq b$ (X bounded random variable). Then

$$\mathbb{E}[e^{sX}] \leq \exp\left\{\frac{s^2(b-a)^2}{8}\right\} \quad \text{for } s > 0$$

Proof. By the convexity of the exp function we have

$$e^{sx} \leq \frac{x-a}{b-a} e^{sb} + \frac{b-x}{b-a} e^{sa} \quad \text{with } a \leq x \leq b$$

Using $\mathbb{E}[X] = 0$ and defining $p := \frac{-a}{b-a}$ we obtain

$$\begin{aligned}
\mathbb{E}[e^{sX}] &\leq \mathbb{E}\left[\frac{X-a}{b-a}e^{sb} + \frac{b-X}{b-a}e^{sa}\right] \\
&\leq \frac{b}{b-a}e^{sa} - \frac{a}{b-a}e^{sb} \\
&= \frac{b-a+a}{b-a}e^{sa} + pe^{sb} \\
&= (1-p)e^{sa} + pe^{sb} \\
&= (1-p)e^{sa} + pe^{s(b-a+a)} \\
&= (1-p)e^{sa} + pe^{s(b-a)}e^{sa} \\
&= (1-p)e^{sa} + pe^{s(b-a)}e^{sa\frac{b-a}{b-a}} \\
&= (1-p)e^{sa} + pe^{s(b-a)}e^{-ps(b-a)}
\end{aligned}$$

Then defining

$$\mu = s(s-a)$$

$$\Phi(\mu) = -p\mu + \ln(1-p + pe^\mu)$$

so we have that the last equality $(1-p)e^{sa} + pe^{s(b-a)}e^{-ps(b-a)} = e^{\Phi(\mu)}$

It is possible to show

$$\Phi'(X) = -p + \frac{p}{p + (1-p)e^{-\mu}}$$

therefore $\Phi(\mu) = \Phi'(0) = 0$, moreover

$$\Phi(\mu) = \frac{p(1-p)e^{-\mu}}{(p + (1-p)p^{-\mu})^4} \leq \frac{1}{4}$$

by Taylor's theorem we have:

$$\Phi(x) = \Phi(0) + \mu\Phi'(0) + \frac{\mu}{2}\Phi''(\sigma) \leq \frac{\mu^2}{8} = \frac{s^2(b-a)^2}{8}$$

with $\sigma \in [0, \mu]$. □

We're now ready for the **Hoeffding's Inequality**

Theorem 2.1.1. *Let $(x_1 \dots x_n)$ be independent random variable such that $x_i \in [a_i, b_i]$ then for any $t > 0$*

$$\mathbb{P}(S_n - \mathbb{E}[S_n] \geq t) \leq e^{-\frac{2t^2}{\sum_{i=1}^n (b_i - a_i)^2}}$$

$$\mathbb{P}(S_n - \mathbb{E}[S_n] \leq -t) \leq e^{-\frac{2t^2}{\sum_{i=1}^n (b_i - a_i)^2}}$$

Proof. Using the *Chernoff's bounding* for sums of random variables 2.1.2 and the precedent lemma 2.1.1 we obtain

$$\mathbb{P}(S_n - \mathbb{E}[S_n] \geq t) \leq e^{-st} \prod_{i=1}^n e^{\frac{s^2(b_i-a_i)^2}{8}} = e^{-st} e^{\frac{s^2}{8} \sum_{i=1}^n (b_i-a_i)^2} = e^{-\frac{2t^2}{\sum_{i=1}^n (b_i-a_i)^2}}$$

where we chose $s = \frac{4t}{\sum_{i=1}^n (b_i-a_i)^2}$ □

This inequality has the same form as the one based on the central limit theorem except that the average variance σ^2 is replaced by the upper bound $\frac{1}{4} \sum_{i=1}^n (b_i - a_i)^2$. Next we will see *Bernstein Inequality* an inequality that take into account also the variance.

Lemma 2.1.2. Assume that $\mathbb{E}[X_i] = 0$ then if for all X_i , $|X_i| \leq c$ (X_i are bounded):

$$\mathbb{E}[e^{sx_i}] \leq \exp \left\{ s^2 \sigma_i^2 \frac{e^{sc} - 1 - sc}{sc} \right\}$$

where $\sigma_i^2 := \mathbb{E}[X_i^2]$

Proof. define $F_i = \sum_{r=2}^{\infty} s^{r-2} \frac{\mathbb{E}[x_i^r]}{r! \sigma_i^2}$.

Since (for Taylor) $e^{sx} = 1 + sx + \sum_{r=2}^{\infty} s^r \frac{x^r}{r!}$ then taking into account $\mathbb{E}[X_i] = 0$

$$\mathbb{E}[s^{sX_i}] = 1 + s\mathbb{E}[X_i] + \sum_{r=2}^{\infty} s^r \frac{\mathbb{E}[x_i^r]}{r!} = 1 + s^2 \sigma_i^2 F_i \leq e^{s^2 \sigma_i^2 F_i}$$

Because we supposed $|X_i| \leq c$ for each index r we have

$$\mathbb{E}[X_i^r] = \mathbb{E}[X_i^{r-2} X_i^2] \leq \mathbb{E}[c^{r-2} X_i^2] = c^{r-2} \sigma_i^2$$

Thus

$$\begin{aligned} F_i &\leq \sum_{r=2}^{\infty} \frac{s^{r-2} c^{r-2} \sigma_i^2}{r! \sigma_i^2} \\ &= \frac{1}{(sc)^2} \sum_{r=2}^{\infty} \frac{(sc)^r}{r!} \\ &= \frac{e^{sc} - 1 - sc}{(sc)^2} \end{aligned}$$

where in the last step we recognized the summation as the exponential wrote in Taylor series missing the first two terms \square

Theorem 2.1.2. Bernstein Inequality

Let $(x_1 \dots x_n)$ be independent real valued random variables with $\mathbb{E}[X_i] = 0$ and $|X_i| \leq c$. Set $\sigma^2 = \frac{1}{n} \sum_{i=1}^n \text{Var}[X_i]$ (note that $\text{Var}[X_i] = \mathbb{E}[X_i^2]$ because $\mathbb{E}[X_i] = 0$). Then for $t > 0$

$$\mathbb{P}\left(\sum_{i=1}^n X_i \geq t\right) \leq \exp \left\{ -\frac{n\sigma^2}{c^2} h\left(\frac{ct}{n\sigma^2}\right) \right\}$$

where $h(\mu) = (1 + \mu) \ln(1 + \mu) - \mu$ for $\mu \geq 0$.

Proof. Using the Chernoff's bounding for sums of random variables 2.1.2 we obtain and the precedent lemma 2.1.2 we obtain

$$\mathbb{P}\left(\sum_{i=1}^n X_i \geq t\right) \leq \exp \left\{ \frac{n\sigma^2(e^{sc} - 1 - sc)}{c^2} - st \right\}$$

and the bound is minimized by $s = \frac{1}{c} \ln \left(1 + \frac{tc}{n\sigma^2} \right)$ \square

Corollary 2.1.1. Referring to the Bernstein Inequality there is a lower bound for h :

$$h(\mu) \geq \frac{\sigma^2}{2 + 2\frac{\mu}{\epsilon}}$$

so for $\epsilon > 0$ the Bernstein Inequality becomes:

$$\mathbb{P}\left(\sum_{i=1}^n X_i \geq t\right) \leq \exp \left\{ -\frac{n\epsilon}{2\sigma^2 + \frac{2}{3}c\epsilon} \right\}$$

This result is extremely useful in hypothesis testing ($\mathbb{P}(T_n > t) = \alpha$) because usually to do the test we have to invert the CDF of T_n . With this result we can instead use the second term of the *Bernstein Inequality* as α and then we can isolate the ϵ to find the small t . Sadly this work only if T_n is a sum of independent random variables which however is the most common situation.

We consider now the problem of deriving inequalities for the Variance of functions of independent random variables.

Lemma 2.1.3. *Let \mathcal{X} be some set and let $g : \mathcal{X}^n \rightarrow \mathbb{R}$ be a measurable function. Define $Z := g(X_1 \dots X_n)$ where $(x_1 \dots x_n)$ are independent random variables in \mathcal{X} and $\mathbb{E}_i Z$ the expected value of Z with respect to X_i that is $\mathbb{E}_i Z = \mathbb{E}[Z | X_1 \dots X_{i-1}, X_{i+1} \dots X_n]$. Then*

$$\text{Var}(Z) \leq \sum_{i=1}^n \mathbb{E}[(Z - \mathbb{E}_i Z)^2]$$

Directly from this lemma follows

Theorem 2.1.3. Efron-Stein Inequality *Let $X'_1 \dots X'_n$ be from an independent copy of $X_1 \dots X_n$ and define $Z'_i = g(X_1 \dots X_{i-1}, X'_i, X_{i+1} \dots X_n)$ then*

$$\text{Var}(Z) \leq \sum_{i=1}^n \mathbb{E}[(Z - Z'_i)^2]$$

when $g(X_1 \dots X_n) = \sum_{i=1}^n X_i$ the inequality becomes an equality.

3 Likelihood Function

The *likelihood function* is a function that contains all the statistical information required to make inference.

Definition 3.0.1. Consider a random sample $(X_1 \dots X_n)$ from $X \sim f_X(X; \theta)$, then the distribution of $(X_1 \dots X_n)$ will be:

$$f_{\underline{X}}(\underline{x}; \theta) = \prod_{i=1}^n f_{X_i}(x_i; \theta)$$

when we see $f_{\underline{X}}(\underline{x}; \theta)$ as a function of θ for fixed \underline{x} we call it **likelihood function**

$$\mathcal{L}(\theta, \underline{x}) = \prod_{i=1}^n f_{X_i}(x_i; \theta)$$

An important function related to the likelihood function is the **log likelihood function**

Definition 3.0.2. The \ln of the likelihood function is said *log likelihood function*

$$V_n(\theta) = \ln \mathcal{L}(\theta, \underline{x}) = \ln \left(\prod_{i=1}^n f_{X_i}(x_i; \theta) \right) = \sum_{i=1}^n \ln(f_{X_i}(x_i; \theta))$$

Definition 3.0.3. Score function:

$$V'_n = \frac{d}{d\theta} V_n(\theta) = \frac{\mathcal{L}'(\theta, \underline{x})}{\mathcal{L}(\theta, \underline{x})}$$

Note that if we fix θ then $\mathcal{L}(\theta, \underline{x})$ is (related to) the probability that the particular value we fixed for θ has generated \underline{x} .

Suppose we fix two value $\theta, \theta_2 \in \Theta$ and $\mathcal{L}(\theta_1, \underline{x}) > \mathcal{L}(\theta_2, \underline{x})$ then we say \underline{x} is "*more likely*" generated under θ_1 .

Note that the same meaning is also applicable to the log likelihood function.

It is because of this that we usually search for the maximum of the likelihood function. Often to find it we just derive, but sometimes \mathcal{L} is not regular enough so we have to "*regularize*" it.

3.1 Likekihood principles

The statistical inference based on the likelihood function is a consequence of two principles:

1. **Week likelihood principle:** for a fixed parametric model $X \sim F_X(x, \theta)$ if two observed samples \underline{x} and \underline{y} are such that

$$\mathcal{L}(\theta, \underline{x}) \propto \mathcal{L}(\theta, \underline{y})$$

then the two likelihood functions are equivalent i.e. must produce the same inference result on θ .

2. **Strong likelihood principle:** let \underline{x} be an observed sample under the model $X \sim F_X(x, \theta)$ with likelihood function $\mathcal{L}(\theta, \underline{x})$ and let \underline{y} be an observed sample under the model $X \sim F_Y(y, \theta)$ with likelihood function $\mathcal{L}(\theta, \underline{y})$.

If $\mathcal{L}(\theta, \underline{x}) \propto \mathcal{L}(\theta, \underline{y})$ the the two samples provides with the same inference.

The fundamental difference between *Probability* and *Statistic* is that in the first one the goal is to find the chance of a random variable to take a particular value, statistic instead given the results of a experiment, try to find the distribution where it came from.

Example 3.1.1. Take (X_1, X_2, X_3) from one of the following distribution:

- (a) $X \sim \text{Ber}(\theta_1), \theta_1 = \frac{1}{2}$
- (b) $X \sim \text{Ber}(\theta_2), \theta_2 = \frac{1}{3}$
- (c) $X \sim \text{Ber}(\theta_3), \theta_3 = \frac{1}{4}$

$X \in [0, 1]$ $\Theta = [0, 1]$.

We can imagine (x_1, x_2, x_3) as the results of a experiment where we had to flip a single coin 3 times. Now we want to know the parameter θ of the coin we flipped tree times and we have tree possibilities: $\theta_1 = \frac{1}{2}, \theta_2 = \frac{1}{3}, \theta_3 = \frac{1}{4}$.

So $(x_1, x_2, x_3) \in \{0, 1\}^3$

| x_1, x_2, x_3 | $\theta_1 = \frac{1}{2}$ | $\theta_2 = \frac{1}{3}$ | $\theta_3 = \frac{1}{4}$ |
|-----------------|--------------------------|--------------------------|--------------------------|
| 0,0,0 | $\frac{1}{8}$ | $\frac{8}{27}$ | $\frac{27}{64}$ |
| 0,0,1 | $\frac{1}{8}$ | $\frac{4}{27}$ | $\frac{9}{64}$ |
| 0,1,0 | $\frac{1}{8}$ | $\frac{4}{27}$ | $\frac{9}{64}$ |
| 1,0,0 | $\frac{1}{8}$ | $\frac{4}{27}$ | $\frac{9}{64}$ |
| 0,1,1 | $\frac{1}{8}$ | $\frac{2}{27}$ | $\frac{3}{64}$ |
| 1,0,1 | $\frac{1}{8}$ | $\frac{2}{27}$ | $\frac{3}{64}$ |
| 1,1,0 | $\frac{1}{8}$ | $\frac{2}{27}$ | $\frac{3}{64}$ |
| 1,1,1 | $\frac{1}{8}$ | $\frac{1}{27}$ | $\frac{1}{64}$ |

Once we know the result of the throw we will "guess" the value of θ choosing the one that give us more probability for the given result.

3.2 Condition of Regularity

In our investigations on θ we will assume some condition of regularity for our model.

Given $X \sim F_x(x, \theta)$

1. we assume that $\theta \in \Theta$, where Θ is a open real set
2. for any $\theta \in \Theta$, there exist the derivative of $\mathcal{L}(\theta; z)$ with respect to θ at least up to the third order
3. for any $\theta_0 \in \Theta$, there exist three functions g, h, H , that are integrable in a neighborhood of θ_0 and

- $\left| \frac{d}{d\theta} f_X(x, \theta) \right| \leq g(x)$
- $\left| \frac{d^2}{d^2\theta} f_X(x, \theta) \right| \leq h(x)$
- $\left| \frac{d^3}{d^3\theta} \ln(f_X(x, \theta)) \right| \leq H(x)$

4. for any $\theta \in \Theta$

$$0 < \mathbb{E}[(\ln(\mathcal{L}(\theta, \underline{X})))^2] < \infty$$

In addition there is the condition of identifiability.

5. We say that a statistical model is identifiable if for every θ_1, θ_2 there is al least one event E such that:

$$\mathbb{P}(X \in E | \theta_1) \neq \mathbb{P}(X \in E | \theta_2)$$

(We will always take 5. as granted)

3.3 Properties of the Likelihood Function

Proposition 3.3.1. *Some properties of the score function 3.0.3 are:*

1. $\mathbb{E}[V'_n(\theta)] = 0$
2. $\text{Var}(V'_n(\theta)) = \mathbb{E}[(V'_n(\theta))^2] = -\mathbb{E}[V''_n(\theta)]$

Proof. 1.

$$\begin{aligned}
 \mathbb{E}[V'_n(\theta)] &= \int_{\mathbb{R}^n} V'_n(\theta) f_{\underline{X}}(\underline{x}, \theta) d\underline{x} \\
 &= \int_{\mathbb{R}^n} \frac{f'_{\underline{X}}(\underline{x}, \theta)}{f_{\underline{X}}(\underline{x}, \theta)} f_{\underline{X}}(\underline{x}, \theta) d\underline{x} \\
 &= \int_{\mathbb{R}^n} \frac{d}{d\theta} f_{\underline{X}}(\underline{x}, \theta) d\underline{x} \\
 &= \frac{d}{d\theta} \int_{\mathbb{R}^n} f_{\underline{X}}(\underline{x}, \theta) d\underline{x} \\
 &= \frac{d}{d\theta} 1 \\
 &= 0
 \end{aligned}$$

Where in the in the fourth equal we used Leibniz and for the fifth recall that $f_{\underline{X}}(\underline{x}, \theta)$ is the PDF of \underline{X}

2. Start by showing that $V''_n(\theta) = \frac{f''_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}; \theta)} - \left(\frac{f'_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}; \theta)} \right)^2$

$$\begin{aligned}
 V''_n(\theta) &= \frac{d^2}{d\theta^2} \ln \mathcal{L}(\theta, \underline{x}) \\
 &= \frac{d}{d\theta} \frac{f'_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}; \theta)} \\
 &= \frac{f''_{\underline{X}}(\underline{x}; \theta) f_{\underline{X}}(\underline{x}; \theta) - f'_{\underline{X}}(\underline{x}; \theta) f'_{\underline{X}}(\underline{x}; \theta)}{|f_{\underline{X}}(\underline{x}; \theta)|^2} \\
 &= \frac{f''_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}; \theta)} - \left(\frac{f'_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}; \theta)} \right)^2 \\
 &= \frac{f''_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}; \theta)} - (V'_n(\theta))^2
 \end{aligned}$$

So now

$$\begin{aligned}
 \mathbb{E}[V''_n(\theta)] &= \int_{\mathbb{R}^n} \frac{f''_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}; \theta)} f_{\underline{X}}(\underline{x}; \theta) d\underline{x} - \mathbb{E}[V'_n(\theta)^2] \\
 &= \int_{\mathbb{R}^n} \frac{d^2}{d\theta^2} f_{\underline{X}}(\underline{x}; \theta) d\underline{x} - \mathbb{E}[V'_n(\theta)^2] \\
 &= \frac{d^2}{d\theta^2} \int_{\mathbb{R}^n} f_{\underline{X}}(\underline{x}; \theta) d\underline{x} - \mathbb{E}[V'_n(\theta)^2] \\
 &= -\mathbb{E}[V'_n(\theta)^2]
 \end{aligned}$$

□

Definition 3.3.1. we define the **Fisher Information** as:

$$\mathcal{I}_n(\theta) = -\mathbb{E}[V''_n(\theta)]$$

Note that this is the definition of Fisher information just for a particular case, there exist a more general one.

The Fisher information has a central role in statistic because it can be shown that for *unbiased estimators* $\tilde{\theta}$ it holds: $Var(\tilde{\theta}) \geq \frac{1}{\mathcal{I}_n(\theta)}$. So if we can find an estimator with $Var(\tilde{\theta}) = \frac{1}{\mathcal{I}_n(\theta)}$ we are sure that it is the one with the lowest variance.

Proposition 3.3.2. Consider $(x_1 \dots x_n)$ from $X \sim F_X(x; \theta)$ regular then

$$\mathcal{I}_n(\theta) = n\mathcal{I}_1(\theta)$$

Example 3.3.1. Consider $(x_1 \dots x_n)$ from $X \sim Ber(\theta)$

$$\mathcal{L}(\theta; \underline{x}) = \prod_{i=1}^n f_{X_i}(x_i; \theta) = \theta^{\sum_{i=1}^n x_i} (1 - \theta)^{n - \sum_{i=1}^n x_i}$$

$$V_n(\theta) = \ln(\mathcal{L}(\theta; \underline{x})) = \ln(\theta) \sum_{i=1}^n x_i + (\ln(1 - \theta)) \left(n - \sum_{i=1}^n x_i \right)$$

$$V'_n(\theta) = \frac{\sum_{i=1}^n x_i}{\theta} - \frac{n - \sum_{i=1}^n x_i}{1 - \theta}$$

$$V''_n(\theta) = \frac{-\sum_{i=1}^n x_i}{\theta^2} - \frac{-\sum_{i=1}^n x_i}{(1 - \theta)^2}$$

$$\begin{aligned} \mathbb{E}[V'_n(\theta)] &= \frac{1}{\theta} \sum_{i=1}^n \mathbb{E}[x_i] - \frac{1}{1 - \theta} \left(n - \sum_{i=1}^n \mathbb{E}[x_i] \right) \\ &= \frac{n\theta}{\theta} - \frac{1}{1 - \theta} (n - n\theta) \\ &= n - n = 0 \end{aligned}$$

$$\begin{aligned} \mathbb{E}[V''_n(\theta)] &= -\frac{1}{\theta} \sum_{i=1}^n \mathbb{E}[X_i] - \frac{1}{1 - \theta} \left(n - \sum_{i=1}^n \mathbb{E}[X_i] \right) \\ &= -\frac{n\theta}{\theta^2} - \frac{n - n\theta}{(1 - \theta)^2} \\ &= -\frac{n}{\theta} - \frac{n(1 - \theta)}{(1 - \theta)^2} \\ &= -\frac{n\theta}{(1 - \theta)\theta} \end{aligned}$$

$$\begin{aligned} \mathcal{I}_n(\theta) &= -\mathbb{E}[V''_n(\theta)] \\ &= \frac{n}{(1 - \theta)\theta} \end{aligned}$$

3.4 Exponential Families

Definition 3.4.1. We say that the distribution of a random variable is an element of an **Exponential Family** $X \sim EF(\theta)$, if its PDF can be written as follows:

$$f_X(x; \theta) = \exp\{Q(\theta)A(x) + C(x) - k(\theta)\}$$

The generalization to random sample becomes:

Definition 3.4.2. We say that the distribution of a random sample $(x_1 \dots x_n)$ from $X \sim EF(\theta)$ is an

element of an **Exponential Family** $X \sim EF(\theta)$ if its PDF can be written as follow:

$$f_{\underline{X}}(\underline{x}; \theta) = \exp\{Q(\theta) \sum_{i=1}^n A(x_i) + \sum_{i=1}^n C(x_i) - nK(\theta)\}$$

Example 3.4.1. $X \sim Ber(\theta)$, $X \in \{0, 1\}$, $\Theta = (0, 1)$

$$\begin{aligned} P_X(x) &= \theta^x (1 - \theta)^{1-x} \mathbb{1}_{\{0,1\}}(x) \\ &= \exp\{x \ln(\theta) + (1 - x) \ln(1 - \theta)\} \mathbb{1}_{\{0,1\}}(x) \\ &= \exp\{x \ln(\theta) + \ln(1 - \theta) - x \ln(1 - \theta)\} \mathbb{1}_{\{0,1\}}(x) \\ &= \exp\left\{x \ln\left(\frac{\theta}{1 - \theta}\right) + \ln(1 - \theta)\right\} \mathbb{1}_{\{0,1\}}(x) \end{aligned}$$

so we get

$$Q(\theta) = \ln\left(\frac{\theta}{1 - \theta}\right)$$

$$A(x) = x$$

$$C(x) = 0$$

$$K(\theta) = -\ln(1 - \theta)$$

Note that for $K(\theta)$ we had to put a $-$ because in the definition we have $-K$.

Proposition 3.4.1. Let $X \sim EF(\theta)$ then

1. $\mathbb{E}[A(X)] = \frac{K'(\theta)}{Q'(\theta)}$
2. $\text{Var}(A(X)) = \frac{K(\theta)}{(Q'(\theta))^2} - \frac{Q''(\theta)}{(Q'(\theta))^2} \frac{K'(\theta)}{Q'(\theta)}$

Note that this proposition gives us only the expectation and variance for $A(X)$, but it is not a problem because usually $A(X) = X$.

Proof. 1. because the exponential family is regular we can use Leibniz so

$$\begin{aligned} 0 &= \frac{d}{d\theta} 1 \\ &= \frac{d}{d\theta} \int f_X(x; \theta) dx \\ &= \int \frac{d}{d\theta} f_X(x; \theta) dx \\ &= \int (A(x)Q'(\theta) - K'(\theta)) f_X(x; \theta) dx \\ &= Q'(\theta) \int A(x) f_X(x; \theta) dx - K'(\theta) \int f_X(x; \theta) dx \\ &= Q'(\theta) \mathbb{E}[A(X)] - K'(\theta) \\ &\implies \mathbb{E}[A(X)] = \frac{K'(\theta)}{Q'(\theta)} \end{aligned}$$

2. because the exponential family is regular we can use Leibniz so

$$\begin{aligned}
0 &= \frac{d^2}{d\theta^2} \int f_X(x; \theta) \\
&= \int \frac{d^2}{d\theta^2} f_X(x; \theta) \\
&= \int (A(x)Q''(\theta) - K''(\theta))f_X(x; \theta) + (A(x)Q'(\theta) - K'(\theta))^2 f_X(x; \theta) dx \\
&= Q''(\theta)\mathbb{E}[A(X)] - K''(\theta) + (Q'(\theta))^2 \int \left(A(x) - \frac{K'(\theta)}{Q'(\theta)} \right)^2 f_X(x; \theta) dx \\
&= Q''(\theta) \frac{K'(\theta)}{Q'(\theta)} - K''(\theta) + (Q'(\theta))^2 \int (A(x) - \mathbb{E}[A(X)])^2 f_X(x; \theta) dx \\
&= Q''(\theta) \frac{K'(\theta)}{Q'(\theta)} - K''(\theta) + (Q'(\theta))^2 \text{Var}(A(X)) \\
\implies \text{Var}(A(X)) &= \frac{K(\theta)}{(Q'(\theta))^2} - \frac{Q''(\theta)}{(Q'(\theta))^2} \frac{K'(\theta)}{Q'(\theta)} = \frac{K(\theta)}{(Q'(\theta))^2} - \frac{Q''(\theta)}{(Q'(\theta))^2} \mathbb{E}[A(X)]
\end{aligned}$$

□

Observation 2. If $Q(\theta) = \theta$ we get

$$\begin{aligned}
\mathbb{E}[A(X)] &= K'(\theta) \\
\text{Var}(A(X)) &= K''(\theta)
\end{aligned}$$

3.5 Natural Exponential Families

Definition 3.5.1. We say that the distribution of a random sample $(x_1 \dots x_n)$ from $X \sim NEF(\theta)$ is an element of a **Natural Exponential Family** $X \sim NEF(\theta)$ if its PDF can be written as follow:

$$f_X(x; \nu) = \exp\{\nu x + C(x) - K(\nu)\}$$

4 Statistics

The notation of statistic was introduced by Fisher (1920).

The importance of sufficiency is that it can be found in any statistical decision (point estimation, testing, confidential bound)

Definition 4.0.1. Let $X = (X_1 \dots X_n)$ be a random sample from a parametric model $X \sim f_X(x, \theta)$ for some $\theta \in \Theta$ unknown.

We say that $T_n = T(X)$ is **sufficient for the parameter θ** if the conditional distribution of X given T_n does not depend of θ i.e. defined:

- $f_{\underline{X}|T_n=t}(\underline{x}; t, \theta)$ the conditional distribution of \underline{X} given T_n
- $h_{\underline{X}, T_n}(z, t, \theta)$ the joint distribution of \underline{X} and T_n
- $g_{T_n}(t, \theta)$ the marginal distribution of T_n

then T_n is **sufficient for the parameter θ** only if $f_{\underline{X}|T_n=t}(\underline{x}; t, \theta)$ does not depend of θ . Note that

$$f_{\underline{X}|T_n=t}(\underline{x}; t, \theta) = \frac{h_{\underline{X}, T_n}(\underline{x}, t, \theta)}{g_{T_n}(t, \theta)}$$

Example 4.0.1. $(X_1 \dots X_n) \in \{0, 1\}^n$ from a $Ber(\theta)$, $\theta \in (0, 1)$. Define $T_n = \sum_{i=1}^n X_i$, then we want to verify if T_n is sufficient.

- $f_{\underline{X}}(\underline{x}; \theta) = \prod_{i=1}^n f_{X_i}(x_i; \theta) = \theta^{\sum_{i=1}^n x_i} (1 - \theta)^{n - \sum_{i=1}^n x_i}$
- $g_{T_n}(t, \theta) = \binom{n}{t} \sigma^t (1 - \sigma)^{n-t} \mathbb{1}_{0,1 \dots n}(t)$
- $h_{\underline{X}, T_n}(z, t, \theta) = \mathbb{P}(\underline{X} = \underline{x}, T_n = t) = \sigma^t (1 - \sigma)^{n-t}$

so

$$f_{\underline{X}|T_n=t} = \frac{\sigma^t (1 - \sigma)^{n-t}}{\binom{n}{t} \sigma^t (1 - \sigma)^{n-t}} = \frac{1}{\binom{n}{t}}$$

So T_n is a sufficient statistic for θ .

This is a really special case because all the X_i are already in function of T_n .

Observation 3. If T_n is sufficient for θ then all the statistical information of θ contained in the random sample is relocated in T_n . In the example above to infer about θ we just need $\sum_{i=1}^n X_i$.

Observation 4. The notation of sufficiency derive from the probability structure of the parametric family $X \sim f_X(x; \theta)$. We can talk about sufficiency for a parameter θ only after we have specified $X \sim f_X(x; \theta)$

The definition of sufficiency based on conditional probability is not of practical use because we need this two distributions $\begin{cases} g_{T_n}(\cdot) \\ h_{\underline{X}, T_n}(\cdot, \cdot) \end{cases}$ that can be difficult to find. To avoid that we could use a corollary of the *Fisher Factorization Theorem*:

Corollary 4.0.1. Let $\underline{X} = (X_1 \dots X_n)$ from $X \sim f_X(x, \theta)$. Then a statistic T_n is sufficient for θ if and only if there exist two non negative functions $g(\cdot), h(\cdot)$ such that $\mathcal{L}(\theta; \underline{x}) = g(T(\underline{x}); \theta) h(\underline{x})$

Observation 5. • g is a function of the observed sample via T_n

• h is a function of the observed sample and does not depend on θ

Example 4.0.2. Recall the example 4.0.1 $(X_1 \dots X_n) \in \{0, 1\}^n$ from a $Ber(\theta)$, $\theta \in (0, 1)$. Define $T_n = \sum_{i=1}^n X_i$, then we want to verify if T_n is sufficient.

We have

$$f_{\underline{X}}(\underline{x}; \theta) = \prod_{i=1}^n f_{X_i}(x_i; \theta) = \theta^{\sum_{i=1}^n x_i} (1 - \theta)^{n - \sum_{i=1}^n x_i}$$

hence we can apply the previous theorem 4.0.1 using

$$1. h(\underline{X}) = 1$$

$$2. g(T_n(\underline{X}); \theta) = \theta^{\sum_{i=1}^n x_i} (1 - \theta)^{n - \sum_{i=1}^n x_i}$$

Example 4.0.3. $(X_1 \dots X_n) \in \{0, 1\}^n$ from a $N(\theta, 1)$. We want to verify that $T_n = \sum_{i=1}^n X_i$ is a sufficient statistic:

$$\begin{aligned} \mathcal{L}(\theta; \underline{x}) &= \prod_{i=1}^n \frac{1}{\sqrt{2\pi}} \exp \left\{ -\frac{1}{2}(x_i - \theta)^2 \right\} \\ &= (2\pi)^{-n/2} \exp \left\{ -\frac{1}{2} \sum_{i=1}^n (x_i - \theta)^2 \right\} \\ &= (2\pi)^{-n/2} \exp \left\{ -\frac{1}{2} \sum_{i=1}^n x_i^2 - \frac{n\theta^2}{2} + \theta \sum_{i=1}^n x_i \right\} \\ &= (2\pi)^{-n/2} \exp \left\{ -\frac{1}{2} \sum_{i=1}^n x_i^2 \right\} \exp \left\{ -\frac{n\theta^2}{2} + \theta \sum_{i=1}^n x_i \right\} \end{aligned}$$

so

$$\bullet h(\underline{x}) = \exp \left\{ -\frac{1}{2} \sum_{i=1}^n x_i^2 \right\}$$

$$\bullet g(\sum_{i=1}^n x_i, \theta) = \exp \left\{ -\frac{n\theta^2}{2} + \theta \sum_{i=1}^n x_i \right\}$$

Theorem 4.0.1. Fisher Theorem

If $f_{\underline{X}}(\underline{x}; \theta)$ is the joint density function or the joint probability mass function of \underline{X} and $q(t; \theta)$ is the density function or the probability mass function of $T_n(\underline{X})$, then $T_n(\underline{X})$ is sufficient for θ if for every point in the sample space, the ratio

$$\frac{f_{\underline{X}}(\underline{x}; \theta)}{q(t; \theta)}$$

is a constant function of θ .

Proof. MISSING □

We can see now the proof of the corollary 4.0.1

Corollary 4.0.2. Savage

Let $f_{\underline{X}}(\underline{x}; \theta)$ be the joint PDF or PMF of a random sample $\underline{X} = (X_1 \dots X_n)$. A statistic T_n is sufficient for θ if and only if there exist two non negative functions $g(t, \theta), h(\underline{x})$ such that for all \underline{x} in the sample space and for all $\theta \in \Theta$

$$f_{\underline{X}}(\underline{x}; \theta) = g(T(\underline{x}); \theta) h(\underline{x})$$

Proof. We are going to prove the theorem only in the discrete settings.

" \Rightarrow " Suppose that $T(\underline{X})$ is sufficient for θ .

Define:

$$\begin{aligned} - g(t, \theta) &:= \mathbb{P}(T(\underline{X}) = t) \\ - h(\underline{x}) &:= \mathbb{P}\left(\underline{X} = \underline{x} \mid T(\underline{X}) = T(\underline{x})\right) \end{aligned}$$

Because $T(\underline{X})$ is sufficient for θ the conditional probability defining $h(\underline{x})$ does not depend on θ . Hence the choice of $g(t, \theta)$ and $h(\underline{x})$ is legitimate and for this choice we have

$$\begin{aligned} \mathbb{P}(\underline{X} = \underline{x}) &= \mathbb{P}(\underline{X} = \underline{x} \wedge T(\underline{X}) = T(\underline{x})) \\ &= \mathbb{P}(T(\underline{X}) = T(\underline{x}))\mathbb{P}(\underline{X} = \underline{x} \mid T(\underline{X}) = T(\underline{x})) \\ &= g(t, \theta)h(\underline{x}) \end{aligned}$$

So we have the factorization and in particular we can see that

$$\mathbb{P}\left(T(\underline{X}) = T(\underline{x})\right) = g(t, \theta)$$

$\Rightarrow g(T(\underline{x}), \theta)$ is the PMF of $T(s)$

" \Leftarrow " We assume that the factorization holds.

Let $q(t, \theta)$ be the PMF of $T(\underline{X})$. We study the ratio

$$\frac{f_{\underline{X}}(\underline{x}; \theta)}{q(T(\underline{x}); \theta)}$$

in particular define

$$A_{T(\underline{x})} = \{\underline{y} \mid T(\underline{y}) = T(\underline{x})\}$$

Then

$$\begin{aligned} \frac{f_{\underline{X}}(\underline{x}; \theta)}{q(T(\underline{x}); \theta)} &= \frac{g(T(\underline{x}); \theta)h(\underline{x})}{q(T(\underline{x}); \theta)} \\ &= \frac{g(T(\underline{x}); \theta)h(\underline{x})}{\sum_{\underline{y} \in A_{T(\underline{x})}} g(T(\underline{x}); \theta)h(\underline{y})} \\ &= \frac{g(T(\underline{x}); \theta)h(\underline{x})}{g(T(\underline{x}); \theta) \sum_{\underline{y} \in A_{T(\underline{x})}} h(\underline{y})} \\ &= \frac{h(\underline{x})}{\sum_{\underline{y} \in A_{T(\underline{x})}} h(\underline{y})} \end{aligned}$$

This is constant with respect to θ .

Then by the Fisher Theorem $T(\underline{X})$ is sufficient for θ .

□

Example 4.0.4. $(X_1 \dots X_n) \in \{0, 1\}^n$ from a $N(\mu, \sigma^2)$, σ^2 known. As we did in the example 4.0.3 we want to find if $T_n = \frac{1}{n} \sum_{i=1}^n X_i$ is a sufficient statistic for μ .

$$\begin{aligned} f_{\underline{X}}(\underline{x}; \mu, \sigma^2) &= (2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right\} \\ &= (2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \bar{x}_n - \bar{x}_n - \mu)^2 \right\} \\ &= (2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \left(\sum_{i=1}^n (x_i - \bar{x}_n)^2 + n(\bar{x}_n - \mu)^2 \right) \right\} \end{aligned}$$

we already know the distribution of $T(\underline{x}) = \bar{x}_n = \frac{1}{n} \sum_{i=1}^n x_i$ is $\bar{X}_n \sim N(\mu, \sigma^2/n)$.
So we can apply Fisher theorem to the ratio:

$$\frac{(2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \left(\sum_{i=1}^n (x_i - \bar{x}_n)^2 + n(\bar{x}_n - \mu)^2 \right) \right\}}{(2\pi\sigma^2/n)^{-1/2} \exp \left\{ -\frac{n}{2\sigma^2} (\bar{x}_n - \mu)^2 \right\}}$$

$$\frac{(2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \bar{x}_n)^2 \right\} \exp \left\{ -\frac{n}{2\sigma^2} (\bar{x}_n - \mu)^2 \right\}}{(2\pi\sigma^2/n)^{-1/2} \exp \left\{ -\frac{n}{2\sigma^2} (\bar{x}_n - \mu)^2 \right\}}$$

$$\frac{(2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \bar{x}_n)^2 \right\}}{(2\pi\sigma^2/n)^{-1/2}}$$

Hence by Fisher Theorem $T_n = \frac{1}{n} \sum_{i=1}^n X_i$ is sufficient for μ

Observation 6. Until now we found only one sufficient statistic for a fixed parametric model. However we can define many sufficient statistics.

For example the statistic given by the identity $T(\underline{x}) = \underline{x}$ is always a sufficient statistic, indeed we can factorize the distribution $f_X(x, \theta)$ with

- $h(x) = 1$
- $g(T(x); \theta) = f_X(x, \theta)$

Observation 7. Given one sufficient statistic a way to produce more sufficient statistics is thru a one to one function.

Suppose $T(\underline{x})$ is a sufficient statistic for θ , and define $T^*(\underline{x}) = r(T(\underline{x}))$ where r is a one to one function with inverse r^{-1} .

By Savage's Theorem there exist g, h such that

$$\mathcal{L}(\theta; \underline{x}) = g(T(\underline{x}), \theta)h(\underline{x}) = g(r^{-1}(r(T(\underline{x})), \theta))h(\underline{x}) = g(r^{-1}(T^*(\underline{x})), \theta)h(\underline{x})$$

So defining $g^*(t, \theta) = g(r^{-1}(t), \theta)$ we have that

$$\mathcal{L}(\theta; \underline{x}) = g^*(T^*(\underline{x}), \theta)h(\underline{x})$$

\implies by Savage Theorem we have that $T^*(\underline{x})$ is a sufficient statistic.

We saw that in principle we can define many sufficient statistics so it is natural to define a tool that allows us to decide when a sufficient statistic is better than another.

Recall that the purpose of statistic is to achieve data reduction without loss of information.

Therefore a statistic that achieve the most data reduction while still retaining all of the information about θ might be preferable.

Observation 8. We saw in example 4.0.3 that if $(X_1 \dots X_n) \in \{0, 1\}^n$ from a $N(\theta, 1)$, $T_n = \sum_{i=1}^n x_i$ is a sufficient statistic. Instead of $\sum_{i=1}^n X_i$ we can use $T'(\underline{x}) \left(\sum_{i=1}^n x_i, \sum_{i=1}^n x_i^2 \right)$. Clearly $T(X)$ is a greater data reduction than $T'(\underline{x})$ since we do not need to know the sample variance if we want to know θ . Moreover we can write $T(\underline{x})$ as a function of $T'(\underline{x})$ by defining the function $r(a, b) = a$, then we can write

$$T(\underline{x}) = \bar{x}_n = r(\bar{x}_n, S_n^2) = r(T'(\underline{x}))$$

Since $T(\underline{x})$ and $T'(\underline{x})$ are both sufficient they contains the same information about μ . In other terms the additional information given by the sample variance is null.

Definition 4.0.2. A sufficient statistic $T(\underline{x})$ is called **minimal** if for any other sufficient statistic $T'(\underline{x})$, $T(\underline{x})$ is a function of $T'(\underline{x})$.

NOTE:

To say that $T(\underline{x})$ is a function of $T'(\underline{x})$ simply means that if $T'(x) = T'(y)$ then $T(x) = T(y)$.

In other terms if $\{B_t\}$ where $B_t := \{t' : T'(t) = T'(t')\}$ is the partition set induced by T' and $\{A_t\}$ where $A_t := \{t' : T(t) = T(t')\}$ is the partition set induced by T then for every t , $B_t \subseteq A_t$.

\implies the partition of the sample space induced by a minimal statistic is the partition with the smallest cardinality.

Theorem 4.0.2. Lehmann and Sheffe let $f_{\underline{X}}(\underline{x}; \theta)$ be the joint density function or joint probability mass function of a random sample $\underline{X} = (X_1 \dots X_n)$. Suppose there exist a function T such that for any two sample points \underline{x} , \underline{y} the ratio

$$\frac{f_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{y}; \theta)}$$

is constant as a function of θ if and only if $T(\underline{x}) = T(\underline{y})$

Then T is a minimal sufficient statistic for θ

Proof. To simplify the proof we assume $f_{\underline{X}}(\underline{x}; \theta) > 0 \forall \underline{x}, \forall \theta$.

First we show that $T(\underline{x})$ is sufficient.

Define \mathcal{T} as the image of the sample space under the function $T(\underline{x})$.

$$\mathcal{T} := \{t : t = T(\underline{x}) \text{ for some } \underline{x} \text{ in the sample space}\}$$

Define $\{A_t\}$ the partition set induced by T , where $A_t := \{\underline{x}' : T(\underline{x}') = T(\underline{x})\}$ For each A_t choose and fix some elements $\underline{x}_t \in A_t$. For any point in the space $\underline{x}_{T(\underline{x})}$ is the fixed element that is in the same set A_t , as \underline{x} . Since \underline{x} and $\underline{x}_{T(\underline{x})}$ are in the same set A_t then $T(\underline{x}) = T(\underline{x}_{T(\underline{x})})$ so by the assumptions the ratio

$$\frac{f_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}_{T(\underline{x})}; \theta)}$$

Does not depend on θ . Thus we can define $h(\underline{x}) := \frac{f_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}_{T(\underline{x})}; \theta)}$.

Then define the function $g(\underline{x}, \theta) f_{\underline{X}}(\underline{x}; \theta)$, so we have:

$$f_{\underline{X}}(\underline{x}; \theta) = \frac{f_{\underline{X}}(\underline{x}_{T(\underline{x})}; \theta) f_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{x}_{T(\underline{x})}; \theta)} = g(T(\underline{x}), \theta) h(\theta)$$

and by Savage Theorem $T(\underline{x})$ is sufficient for θ .

Now we will show that $T(\underline{x})$ is minimal sufficient.

Let $T'(\underline{x})$ be another sufficient statistic. By Savage Theorem we know that exist h', g' such that

$$f_{\underline{X}}(\underline{x}; \theta) = g'(T'(\underline{x}), \theta) h'(\theta)$$

Let $\underline{x}, \underline{y}$ be two sample points such that $T'(\underline{x}) = T'(\underline{y})$ then we can study the ratio:

$$\frac{f_{\underline{X}}(\underline{x}; \theta)}{f_{\underline{X}}(\underline{y}; \theta)} = \frac{g'(T'(\underline{x}); \theta) h'(\theta)}{g'(T'(\underline{y}); \theta) h'(\theta)} = \frac{h'(\underline{x})}{h'(\underline{y})}$$

Since the ratio does not depend on θ , by the assumption (the other implication of the IFF) implies $T(\underline{x}) = T(\underline{y})$. So we can say that $T(\underline{x})$ is a function of $T'(\underline{x})$ therefore $T(\underline{x})$ is minimal. \square

4.1 Estimators

Definition 4.1.1. Suppose there is a fixed parameter θ that needs to be estimated. Then an **estimator** is a function that maps the sample space to a set of sample estimates. An estimator of θ is usually denoted by the symbol $\hat{\theta}$.

Now we are going to introduce some definition of "good" estimators.

Definition 4.1.2. $T_n(\underline{X})$ is said to be **unbiased** for θ if $\mathbb{E}[T_n(\underline{X})] = \theta$

Observation 9. we use the expected value to define a "good" estimator because of the linearity of the operator.

Definition 4.1.3. Bias:

$$Bias_{\theta}(T_n(\underline{X})) = \mathbb{E}\left[T_n(\underline{X}) - \mathbb{E}[T_n(\underline{X})]\right]$$

When we ask an estimator to be unbiased basically we are requiring it to be centred around θ . Another parameter that gives us information about the goodness of an estimator is the variance. We can interpret the variance as a measure of the dispersion around the expected value, so before checking the variance we must be sure that the expected value overlaps with θ . In this scenario the less is the variance the better is the estimator.

Observation 10. Variance is a good parameter to watch only if the estimator is unbiased

To avoid this problem we can introduce the *Mean Squared Error*

Definition 4.1.4. Mean Squared Error(MSE)

$$\mathbb{E}[(T_n(\underline{X}) - \theta)^2]$$

The importance of this quantity comes from the *Chebyshev's Inequality* 1.1.1

$$\mathbb{P}(|T_n(\underline{x}) - \theta| < k) > 1 - \frac{\mathbb{E}[(T_n(\underline{X}) - \theta)^2]}{k^2}$$

Indeed we notice the smaller the MSE the greater is $\mathbb{P}(|T_n(\underline{x}) - \theta| < k)$.

Proposition 4.1.1.

$$\mathbb{E}[(T_n(\underline{X}) - \theta)^2] = Var(T_n(\underline{X})) + Bias_{\theta}(T_n(\underline{X}))^2$$

Proof.

$$\begin{aligned} \mathbb{E}[(T_n(\underline{X}) - \theta)^2] &= \mathbb{E}[(T_n(\underline{X}) - \mathbb{E}[T_n(\underline{X})] + \mathbb{E}[T_n(\underline{X})] - \theta)^2] \\ &= \mathbb{E}[(T_n(\underline{X}) - \mathbb{E}[T_n(\underline{X})])^2] + \mathbb{E}[(\mathbb{E}[T_n(\underline{X})] - \theta)^2] + 2\mathbb{E}[(T_n(\underline{X}) - \mathbb{E}[T_n(\underline{X})])(\mathbb{E}[T_n(\underline{X})] - \theta)] \\ &= \mathbb{E}[(T_n(\underline{X}) - \mathbb{E}[T_n(\underline{X})])^2] + \mathbb{E}[(\mathbb{E}[T_n(\underline{X})] - \theta)^2] + 2\mathbb{E}[T_n(\underline{X})\mathbb{E}[T_n(\underline{X})] - T_n(\underline{X})\theta - \mathbb{E}[T_n(\underline{X})]\theta] \\ &= \mathbb{E}[(T_n(\underline{X}) - \mathbb{E}[T_n(\underline{X})])^2] + \mathbb{E}[(\mathbb{E}[T_n(\underline{X})] - \theta)^2] \\ &= Var(T_n(\underline{X})) + Bias_{\theta}(T_n(\underline{X}))^2 \end{aligned}$$

□

Observation 11. If $\mathbb{E}[T_n(\underline{X})] = \theta$ then $MSE(T_n(\underline{X})) = Var(T_n(\underline{X}))$.

Definition 4.1.5. Let $X_1..X_n$ from $X \sim f_X(x, \theta)$, T'_n and T''_n estimators for θ . We say T'_n is **more efficient** than T''_n if

$$MSE(T'_n) < MSE(T''_n)$$

Usually we choose the estimator with the lower MSE even if it is biased.

4.2 Properties of Estimators

The problem of the MSE is that we can not be sure that there exist T_n such that $MSE(T_n)$ is the lowest possible.

A solution for this comes from

Theorem 4.2.1. Cramer-Rao Bound

Let $X = (X_1 \dots X_n)$ be a random sample from a parametric model $X \sim f_X(x, \theta)$.

Then, under condition of regularity, for any estimator T_n of θ

$$\text{Var}(T_n) \geq \frac{[1 + b'(T_n)]^2}{\mathcal{I}_n(\theta)}$$

Where

$b(T_n)$ is the bias of T_n

$\mathcal{I}_n(\theta)$ is the Fisher Information

Proof. consider the estimator T_n .

- $\mathbb{E}[T_n] = \theta + b(\theta)$
- $\frac{d}{d\theta} \mathbb{E}[T_n] = 1 + b'(\theta)$
- $\mathbb{E}[V'_n(\theta)] = 0 \leftarrow$ because we suppose our model regular

$$\implies \text{Cov}(T_n, V'_n(\theta)) = \mathbb{E}[T_n V'_n(\theta)] - \mathbb{E}[T_n] \mathbb{E}[V'_n(\theta)] = \mathbb{E}[T_n V'_n(\theta)]$$

$$\begin{aligned} \mathbb{E}[T_n V'_n(\theta)] &= \int_{\mathbb{R}^n} T_n V'_n(\theta) f_X(\underline{x}; \theta) d\underline{x} \\ &= \int_{\mathbb{R}^n} T_n \frac{f'_X(\underline{x}, \theta)}{f_X(\underline{x}; \theta)} f_X(\underline{x}; \theta) d\underline{x} \\ &= \int_{\mathbb{R}^n} T_n \frac{d}{d\theta} f_X(\underline{x}; \theta) d\underline{x} \\ &= \frac{d}{d\theta} \int_{\mathbb{R}^n} T_n f_X(\underline{x}; \theta) d\underline{x} \\ &= \frac{d}{d\theta} \mathbb{E}[T_n] \\ &= 1 + b'(\theta) \end{aligned}$$

So

$$\text{Cov}(T_n, V'_n(\theta)) = 1 + b'(\theta)$$

We know that in general for X, Y random variables such that $\mathbb{E}[X] = \mu, \mathbb{E}[Y] = \nu, \mathbb{E}[X^2] < \infty, \mathbb{E}[Y^2] < \infty$, from Cauchy-Schwarz, it holds

$$\begin{aligned} (\text{Cov}(X, Y))^2 &= \left(\mathbb{E}[(X - \mu)(Y - \nu)] \right)^2 \\ &\leq \mathbb{E}[(X - \mu)^2] \mathbb{E}[(Y - \nu)^2] \\ &= \text{Var}(X) \text{Var}(Y) \end{aligned}$$

So replacing X with T_n and Y with $V'_n(\theta)$ we obtain

$$\begin{aligned} \text{Var}(T_n) &\geq \frac{(\text{Cov}(T_n, V'_n(\theta)))^2}{\text{Var}(V'_n(\theta))} \\ &= \frac{(\text{Cov}(T_n, V'_n(\theta)))^2}{\mathbb{E}[(V'_n(\theta) - \mathbb{E}[V'_n(\theta)])^2]} \\ &= \frac{(\text{Cov}(T_n, V'_n(\theta)))^2}{\mathbb{E}[V'_n(\theta)^2]} \\ &= \frac{(\text{Cov}(T_n, V'_n(\theta)))^2}{\mathcal{I}_n(\theta)} \end{aligned}$$

□

Corollary 4.2.1. *Under condition of regularity*

$$MSE \geq \frac{(1 + b'(\theta))^2}{\mathcal{I}_n(\theta)} + b^2(\theta)$$

Proof. Directly from Cramer-Rao Bound 4.2.1 remembering that $MSE(T_n) = Var(T_n) + b^2(\theta)$ \square

Corollary 4.2.2. *Let $\underline{X} = (X_1 \dots X_n)$ be a random sample from a regular model $X \sim f_X(x; \theta)$. If there exist a unbiased estimator for θ whose variance is equal to the Cramer-Rao bound, Then T_n is unique*

Proof. Take T_{1n}, T_{2n} be unbiased estimators for θ such that

$$Var(T_{1n}) = Var(T_{2n}) = \frac{1}{\mathcal{I}_n(\theta)}$$

Define $T_n := \frac{T_{1n} + T_{2n}}{2}$
 $\mathbb{E}[T_n] = \frac{\mathbb{E}[T_{1n}] + \mathbb{E}[T_{2n}]}{2} = \frac{2}{2}\theta = \theta$
 $\implies T_n$ is also unbiased for θ
 $\implies Var(T_n) \geq \frac{1}{\mathcal{I}_n(\theta)}$

$$\begin{aligned} Var(T_n) &= Var\left(\frac{T_{1n} + T_{2n}}{2}\right) \\ &= \frac{1}{4} \left[Var(T_{1n}) + Var(T_{2n}) + 2Cov(T_{1n}, T_{2n}) \right] \\ &= \frac{1}{4} \left[Var(T_{1n}) + Var(T_{2n}) + 2Cov(T_{1n}, T_{2n}) \right] \frac{[Var(T_{1n})Var(T_{2n})]^{1/2}}{[Var(T_{1n})Var(T_{2n})]^{1/2}} \\ &= \frac{(1 + Corr(T_{1n}, T_{2n}))}{2} \frac{1}{\mathcal{I}_n(\theta)} \end{aligned}$$

MEMO: $|Corr(X, Y)| \leq 1$

Because $Var(T_n) \leq \frac{1}{\mathcal{I}_n(\theta)}$ then we must have $Corr(T_{1n}, T_{2n}) \geq 1 \implies Corr(T_{1n}, T_{2n}) = 1$.

$\implies T_{2n} = a + bT_{1n}$.

Hence we must have $\theta = \mathbb{E}[T_{1n}] = \mathbb{E}[a + bT_{1n}] = \mathbb{E}[a] + b\mathbb{E}[T_{1n}] = a + b\theta \implies a = 0, b = 1$

$\implies T_{2n} = T_{1n}$ \square

Definition 4.2.1. Consider a regular model $X \sim f_x(x\theta)$. We say that an unbiased estimator T_n whose variance is

$$Var(T_n) = \frac{1}{\mathcal{I}_n(\theta)}$$

is **efficient**.

Moreover we define the **efficiency** of an estimator as:

$$Eff(T_n) = \frac{1}{Var(T_n)\mathcal{I}_n(\theta)} \in [0, 1]$$

Observation 12. 1. We introduce (absolute) efficiency at the cost of assuming regularity for the parametric model.

2. The variance of an unbiased estimator T_n of θ can not be smaller than the Cramer-Rao bound. However we do not know if there exist an estimator whose variance is equal to the Cramer-Rao bound.

3. The proper lower bound involves the MSE

$$MSE(T_n) \geq \frac{[1 + b'(\theta)]}{\mathcal{I}_n(\theta)} + b^2(\theta)$$

Proposition 4.2.1. Let $\underline{X} = (X_1 \dots X_n)$ be a random sample from a parametric model $X \sim f_x(x\theta)$. Let T_n be a unbiased estimator for θ . Then T_n is efficient for θ if and only if

$$V'_n(\theta) = \mathcal{I}_n(\theta)(T_n - \theta)$$

Proof.

$$\begin{aligned} V'_n(\theta) = \mathcal{I}_n(\theta)(T_n - \mathbb{E}[T_n]) &\iff V'_n(\theta)^2 = \mathcal{I}_n(\theta)^2(T_n - \mathbb{E}[T_n])^2 \\ &\iff \mathbb{E}[V'_n(\theta)^2] = \mathbb{E}[\mathcal{I}_n(\theta)^2(T_n - \mathbb{E}[T_n])^2] \\ &\iff \mathcal{I}_n(\theta) = \mathcal{I}_n(\theta)^2 \mathbb{E}[(T_n - \mathbb{E}[T_n])^2] \\ &\iff 1 = \mathcal{I}_n(\theta) \text{Var}(T_n) \\ &\iff \text{Var}(T_n) = \frac{1}{\mathcal{I}_n(\theta)} \end{aligned}$$

i.e. T_n is efficient. □

Theorem 4.2.2. Rao-Blackwell

Let $\underline{X} = (X_1 \dots X_n)$ a random sample from a parametric model $X \sim f_x(x\theta)$. Let

- T_{1n} a sufficient estimator for θ
- T_{2n} a unbiased estimator for θ
- $T_n := \mathbb{E}[T_{2n}|T_{1n}]$

Then

1. T_n is a function of T_{1n}
2. $\mathbb{E}[T_n] = \theta$
3. $\text{Var}(T_n) < \text{Var}(T_{2n})$

Definition 4.2.2. Let X_1, X_2, \dots real valued random variables with CDF F_{X_1}, F_{X_2}, \dots . We say that $(X_n)_n$ **converges in distribution** or **converges weakly** to a random variable X if

$$\lim_{n \rightarrow \infty} F_{X_n}(x) = F_X(x)$$

Definition 4.2.3. A sequence $\{X_n\}_n$ of random variables **converges in probability** towards the random variable X if for all $\epsilon > 0$

$$\lim_{n \rightarrow \infty} \mathbb{P}(|X_n - X| > \epsilon) = 0$$

We write $X_n \xrightarrow{p} X$

Definition 4.2.4. Given a real number $r \geq 1$, we say that the sequence $\{X_n\}$ converges in the **r-th mean** (or **in the L^r -norm**) towards the random variable X , if the r-th absolute moments $\mathbb{E}(|X_n|^r)$ and $\mathbb{E}(|X|^r)$ of X_n and X exist, and

$$\lim_{n \rightarrow \infty} \mathbb{E}[|X_n - X|^r] = 0$$

We write $X_n \xrightarrow{L^r} X$.

For $r = 2$ we say that $\{X_n\}$ converges in **mean square** to X .

Proposition 4.2.2. *Convergence in probability \implies convergence in distribution.*

If X is a degenerate random variable we have also

convergence in distribution \implies convergence in probability.

Definition 4.2.5. We say an estimator T_n is **consistent in mean squared** for θ if

$$\lim_{n \rightarrow \infty} MSE(T_n) = 0$$

NOTATION: we will use $b(T_n) := b(\theta)$

Observation 13. Since $MSE(T_n) = Var(T_n) + b^2(T_n)$, then $\lim_{n \rightarrow \infty} MSE(T_n) = 0$ is equivalent to say

- $\lim_{n \rightarrow \infty} Var(T_n) = 0$
- $\lim_{n \rightarrow \infty} b^2(T_n) = 0$

Definition 4.2.6. T_n is **asymptotically unbiased** for θ if

- $\lim_{n \rightarrow \infty} \mathbb{E}[T_n] = \theta$
- $\lim_{n \rightarrow \infty} b(T_n) = 0$

Proposition 4.2.3. A consistent estimator in mean square is also asymptotically unbiased

Definition 4.2.7. We say T_n is consistent in probability for θ if for all $\epsilon > 0$

$$\lim_{n \rightarrow \infty} \mathbb{P}(|T_n - \theta| < \epsilon) = 1$$

Proposition 4.2.4. The consistency in mean squared implies consistency in probability

Proof. Using Chebyshev's inequality 1.1.1

$$\begin{aligned} \mathbb{P}(|T_n - \theta| < \epsilon) &\geq 1 - \frac{MSE(T_n)}{\epsilon^2} \\ \lim_{n \rightarrow \infty} \mathbb{P}(|T_n - \theta| < \epsilon) &\geq 1 - \lim_{n \rightarrow \infty} \frac{MSE(T_n)}{\epsilon^2} \end{aligned}$$

□

Theorem 4.2.3. Central Limit Theorem

Let $\underline{X} = (X_1 \dots X_n)$ be a random sample of size n with X_i independent and identically distributed random variables. With expected value μ and finite variance σ^2 . Then $S_n := \frac{\sum_{i=1}^n X_i}{n}$ converges in probability to the expected value μ

$$S_n \xrightarrow{p} \mu$$

Theorem 4.2.4. Weak law of large numbers

Let $\underline{X} = (X_1 \dots X_n)$ be a random sample of size n with X_i independent and identically distributed random variables. With expected value μ . Define $\bar{X}_n = \frac{1}{n} \sum_{i=1}^n X_i$, then

$$\bar{X}_n \xrightarrow{p} \mu$$

That is, for any $\epsilon > 0$

$$\lim_{n \rightarrow \infty} \mathbb{P}(|\bar{X}_n - \mu| > \epsilon) = 0$$

4.3 Method of Moments

Observation 14. Before we start, like the likelihood method, this method can also be done for irregular models.

Let $X = (X_1, \dots, X_n)$ be a random sample from a parametric model $X \sim f_X(x, \theta)$. Assume that $\theta \in \Theta$ s.t. Θ has dimension r . In order to apply the method of moment for estimating θ we need:

- A condition related to the dimension of Θ . If the number of parameters that we need to estimate is r , then the moments of X must exist up to the order r .

$$\mathbb{E}[|x|^r] < \infty.$$

- Explicit expression for the first r moments of X .

Given that, the method of moments exists in solving the system of equations given by:

$$\frac{1}{n} = \sum_{i=1}^n X_i^J = \mathbb{E}[X^J], \quad \forall J = 1, \dots, r$$

Observation 15. We don't have any result on the theoretical qualities on this method for the simulation. We need to check case by case.

4.4 Maximum Likelihood Approach

We saw (exercice week not) that the Maximum Likelihood Approach consist in finding the

$$\arg \max_{\theta \in \Theta} \ln(\mathcal{L}(\theta; \underline{x}))$$

- if the model is regular then we can simply solve the ML equation $V'_n(\theta) = 0$
- if the parameter is a positive integer then we used the ratio
- the parameter defines the support of the model we must insert the indicator function in the likelihood function
- if the model is regular but we can not solve explicitly the likelihood equation.

A way to solve this is the **Newton-Raphson method**

Theorem 4.4.1. Delta Method (Generalization of Central Limit Theorem for functions)
Let $(Y_n)_{n \geq 1}$ be a sequence of random variables such that

$$\sqrt{n}(Y_n - \theta) \xrightarrow{w} N(0, \sigma^2) \quad \text{as } n \rightarrow \infty$$

for a specific function g and a given value θ assume $g'(\theta)$ exist and it is not zero then

$$\sqrt{n}(g(Y_n) - g(\theta)) \xrightarrow{w} N(0, \sigma^2[g'(\theta)]^2) \quad \text{as } n \rightarrow \infty$$

Proof. Take Taylor's approx of $g(Y_n)$ around $Y_0 = \theta$.

Then

$$g(Y_n) = g(\theta) + g'(\theta)[Y_n - \theta] + R$$

Where $R \rightarrow 0$ as $Y_n \rightarrow \theta$. Since $Y_n \xrightarrow{p} \theta$ then $R \xrightarrow{p} 0$. By an application of Slutsky Theorem we have:

$$\sqrt{n}(g(Y_n) - g(\theta)) = \sqrt{n}(Y_n - \theta)g'(\theta) \xrightarrow{w} N(0, \sigma^2[g'(\theta)]^2) \quad \text{as } n \rightarrow \infty$$

□

A useful property of the maximum likelihood estimator is the **invariance property**.

Suppose we have a parametric family indexed by a parameter θ but we are interested in estimating $\mathcal{T}(\theta)$. The invariance property of maximum likelihood estimator says that if $\hat{\theta}$ is the maximum likelihood estimator for θ then $\mathcal{T}(\hat{\theta})$ is the maximum likelihood estimator for $\mathcal{T}(\theta)$.

If the map $\theta \mapsto \mathcal{T}(\theta)$ is one to one then the invariance property is easy to show. Indeed if we let $\mu = \mathcal{T}(\theta)$ then the inverse function $\mathcal{T}^{-1}(\mu) = \theta$ is well defined and the likelihood function written as a function of μ is:

$$\mathcal{L}^*(\mu; \underline{x}) = \prod_{i=1}^n f_{X_i}(x_i; \mathcal{T}^{-1}(\mu)) = \mathcal{L}(\mathcal{T}^{-1}(\mu); \underline{x})$$

and

$$\sup_{\tau} \mathcal{L}^*(\mu; \underline{x}) = \sup_{\tau} \mathcal{L}(\mathcal{T}^{-1}(\mu); \underline{x}) = \sup_{\theta} \mathcal{L}(\theta; \underline{x})$$

Then the max of $\mathcal{L}^*(\mu; \underline{x})$ is attained at $\mathcal{T} = \mathcal{T}(\theta) = \mathcal{T}(\hat{\theta})$. This shows $\mathcal{T}(\hat{\theta})$ is the maximum likelihood estimator for $\mathcal{T}(\theta)$.

In general the invariance property of maximum likelihood estimator is true for any functional of θ .

Theorem 4.4.2. invariance property of maximum likelihood estimator

If $\hat{\theta}$ is the maximum likelihood estimator for θ then for all function $\mathcal{T}(\theta)$ the maximum likelihood estimator of $\mathcal{T}(\theta)$ is $\mathcal{T}(\hat{\theta})$.

Proof. Define

$$\mathcal{L}^*(\mu; \underline{x}) := \sup_{\{\theta \in \Theta: \mathcal{T}(\theta) = \mu\}} \mathcal{L}(\theta; \underline{x})$$

let $\hat{\mu}$ the value that maximize the function $\mathcal{L}^*(\mu; \underline{x})$, it can be shown that $\hat{\mu}$ coincides with the argument that maximize $\mathcal{L}(\theta; \underline{x})$.

We must show that

$$\mathcal{L}^*(\hat{\mu}; \underline{x}) = \mathcal{L}^*(\mathcal{T}(\hat{\theta}); \underline{x})$$

As we stated above $\hat{\mu}$ coincides with the argument that maximize $\mathcal{L}(\theta; \underline{x})$ so

$$\mathcal{L}^*(\hat{\mu}; \underline{x}) = \sup_{\mu} \sup_{\{\theta \in \Theta: \mathcal{T}(\theta) = \mu\}} \mathcal{L}(\theta; \underline{x}) = \sup_{\theta} \mathcal{L}(\theta; \underline{x}) = \mathcal{L}^*(\hat{\theta}; \underline{x})$$

Moreover

$$\begin{aligned} \mathcal{L}^*(\hat{\theta}; \underline{x}) &= \sup_{\{\theta \in \Theta: \mathcal{T}(\theta) = \mathcal{T}(\hat{\theta})\}} \mathcal{L}(\theta; \underline{x}) \\ &= \mathcal{L}^*(\mathcal{T}(\hat{\theta}); \underline{x}) \end{aligned}$$

By combining the two identities we have that

$$\mathcal{L}^*(\hat{\mu}; \underline{x}) = \mathcal{L}^*(\mathcal{T}(\hat{\theta}); \underline{x})$$

So $\mathcal{T}(\hat{\theta})$ is the maximum likelihood estimator of $\mathcal{T}(\theta)$ where $\hat{\theta}$ is the maximum likelihood estimator of θ □

Observation 16. Fore example we can use the above theorem 4.4.2 to compute the maximum likelihood estimator of θ^2 where θ is the mean of a Gaussian model

Observation 17. The above result 4.4.2 holds also in the context of multidimensional parametric spaces.

Example 4.4.1. $\underline{X} = (X_1 \dots X_n)$ from $X \sim N(\mu, \sigma^2)$. Find the maximum likelihood estimator estimator of (μ, σ^2) .

- $\mathcal{L}(\mu, \sigma^2; \underline{x}) = \prod_{i=1}^n \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(x_i - \mu)^2}{2\sigma^2} \right\} = (2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right\}$
- $\ln(\mathcal{L}(\mu, \sigma^2; \underline{x})) = -\frac{n}{2} \ln(2\pi\sigma^2) - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \propto -\frac{n}{2} \ln \sigma^2 - \frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2$
- $\frac{d}{d\mu} \ln(\mathcal{L}(\mu, \sigma^2; \underline{x})) = \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu)$
- $\frac{d}{d\sigma^2} \ln(\mathcal{L}(\mu, \sigma^2; \underline{x})) = -\frac{n}{2\sigma^2} + \frac{1}{2(\sigma^2)^2} \sum_{i=1}^n (x_i - \mu)^2$

so solving the system

$$\begin{cases} \frac{d}{d\mu} \ln(\mathcal{L}(\mu, \sigma^2; \underline{x})) = \frac{1}{\sigma^2} \sum_{i=1}^n (x_i - \mu) = 0 \\ \frac{d}{d\sigma^2} \ln(\mathcal{L}(\mu, \sigma^2; \underline{x})) = -\frac{n}{2\sigma^2} + \frac{1}{2(\sigma^2)^2} \sum_{i=1}^n (x_i - \mu)^2 = 0 \end{cases}$$

we obtain $\hat{\mu} = \frac{1}{n} \sum_{i=1}^n x_i$; $\hat{\sigma}^2 = \frac{1}{n} \sum_{i=1}^n (x_i - \bar{x}_n)^2$.

The question now is whether $(\hat{\mu}, \hat{\sigma}^2)$ is a point of maximum or not. To do this we can check the Jacobian, but the computational work would be too heavy.

Another way is by reducing the dimension of the problem:

Note that

$$\sum_{i=1}^n (x_i - \mu)^2 = \sum_{i=1}^n (x_i - \bar{x}_n + \bar{x}_n - \mu)^2 = \sum_{i=1}^n (x_i - \bar{x}_n)^2 + \sum_{i=1}^n (\bar{x}_n - \mu)^2 + 2 \sum_{i=1}^n (x_i - \bar{x}_n)(\bar{x}_n - \mu) = \sum_{i=1}^n (x_i - \bar{x}_n)^2 + \sum_{i=1}^n (\bar{x}_n - \mu)^2$$

So, for any $\theta \neq \bar{x}_n$ we have:

$$\sum_{i=1}^n (x_i - \mu)^2 > \sum_{i=1}^n (x_i - \bar{x}_n)^2$$

Hence:

$$(2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \bar{x}_n)^2 \right\} \geq (2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n (x_i - \mu)^2 \right\}$$

Therefore the problem of verifying that $(\hat{\mu}, \hat{\sigma}^2)$ is a maximum is reduced to a one dimensional problem (we need to check only $\hat{\sigma}^2$ is actually the max).

Observation 18. In the end we can say that the invariance property is a good property from a computable point of view, but this property is also the reason why the maximum likelihood estimator can be biased for a finite sample size.

Proposition 4.4.1. Let $\underline{X} = (X_1 \dots X_n)$ be a random sample from a parametric model $X \sim f_X(x, \theta)$. If the maximum likelihood estimator T_n for θ exist and it is unique then T_n is a function of the sufficient statistic for θ

Proof. (informal)

If a sufficient statistic exist then the likelihood function $\mathcal{L}(\theta; \underline{x})$ can be factorized as follows:

$$\mathcal{L}(\theta; \underline{x}) = g(T(\underline{x}); \theta) h(\underline{x})$$

$$\implies \ln \mathcal{L}(\theta; \underline{x}) = \ln(g(T(\underline{x}); \theta)) + \ln(h(\underline{x})) \propto \ln(g(T(\underline{x}); \theta))$$

□

Theorem 4.4.3. Let $\underline{X} = (X_1 \dots X_n)$ be a random sample from a parametric model $X \sim f_X(x; \theta)$ and let $\mathcal{L}(\theta; \underline{x})$ be the likelihood function. Let $\hat{\theta}_n$ be the maximum likelihood estimator for θ and let $\mathcal{T}(\theta)$ be a continuous function of θ . Under the assumption of regularity, for any $\epsilon > 0$ and for any point $\theta \in \Theta$ we have that:

$$\lim_{n \rightarrow \infty} \mathbb{P}(|\mathcal{T}(\hat{\theta}_n) - \mathcal{T}(\theta)| \geq \epsilon) = 0$$

I.e. the maximum likelihood estimator are consistent in probability

Theorem 4.4.4. Let $\underline{X} = (X_1 \dots X_n)$ be a random sample from a parametric model $X \sim f_X(x; \theta)$ and let $\mathcal{L}(\theta; \underline{x})$ be the likelihood function. Let $\hat{\theta}_n$ be the maximum likelihood estimator for θ and let $\mathcal{T}(\theta)$ be a continuous function of θ . Under the assumption of regularity, for any $\epsilon > 0$ and for any point $\theta \in \Theta$ we have that:

$$\sqrt{n}[\mathcal{T}(\hat{\theta}_n) - \mathcal{T}(\theta)] \xrightarrow{w} N\left(0, \frac{1}{\mathcal{I}_1(\theta)}\right)$$

where $\mathcal{I}_1(\theta)$ is the Fisher information of the parametric model.

This means that the maximum likelihood estimator are:

- consistent (weakly)

- asymptotically unbiased
- asymptotically Gaussian
- asymptotically efficient

Proof. we will prove the result only for $\mathcal{T}(\theta) = \theta$. The proof for a generic $\mathcal{T}(\theta)$ is an application of the delta method.

We have that $L(\theta; \underline{x}) = \sum_{i=1}^n \ln(f_{X_i}(x_i; \theta))$.

We denote by L' and L'' the first and the second derivative with respect to θ .

Expand the first derivative (Taylor) around θ_0

$$L'(\theta; \underline{x}) = L'(\theta_0; \underline{x}) + (\theta - \theta_0)L''(\theta_0; \underline{x}) + R$$

where R is the residual and can be ignored.?

Now substitute θ with the maximum likelihood estimator of θ , say $\hat{\theta}_n$

$$0 = L'(\hat{\theta}_n; \underline{x}) + (\hat{\theta}_n - \theta_0)L''(\theta_0; \underline{x}) + R$$

$$\begin{aligned} \implies \sqrt{n}(\hat{\theta}_n - \theta_0) &= -\sqrt{n} \frac{L'(\theta_0)}{L''(\theta_0; \underline{x})} \\ &= \frac{-\frac{1}{\sqrt{n}}L'(\theta_0)}{\frac{1}{n}L''(\theta_0; \underline{x})} \end{aligned}$$

We can apply the Central Limit Theorem to the denominator and the Weak Law of Large Numbers to the numerator:

Numerator:

$$\frac{1}{\sqrt{n}}L'(\theta_0; \underline{x}) = \sqrt{n} \left(\frac{1}{n} \sum_{i=1}^n W_i \right)$$

Where $W_i = \frac{\frac{d}{d\theta} f_{X_i}(x_i; \theta)}{f_{X_i}(x_i; \theta)}$ and $\mathbb{E}[W_i] = 0$ and $Var(W_i) = \mathcal{I}_1(\theta)$. By the Central Limit Theorem

$$\frac{1}{\sqrt{n}}L'(\theta_0; \underline{x}) \xrightarrow{w} N(0; \mathcal{I}_1(\theta))$$

Denominator:

$$-\frac{1}{n}L''(\theta_0; \underline{x}) = \frac{1}{n} \sum_{i=1}^n W_i^2 - \frac{1}{n} \sum_{i=1}^n \frac{\frac{d^2}{d\theta^2} f_{X_i}(x_i; \theta)}{f_{X_i}(x_i; \theta)}$$

Where $\mathbb{E}[\frac{1}{n} \sum_{i=1}^n W_i^2] = \mathcal{I}_1(\theta_0)$ and $\mathbb{E}\left[\frac{1}{n} \sum_{i=1}^n \frac{\frac{d^2}{d\theta^2} f_{X_i}(x_i; \theta)}{f_{X_i}(x_i; \theta)}\right] = 0$. Then by the Weak Law of Large Numbers

$$-\frac{1}{n}L''(\theta; \underline{x}) \xrightarrow{w} \mathcal{I}_1(\theta_0)$$

So

$$\sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow{w} \frac{X}{\mathcal{I}_1(\theta_0)}$$

Where $X \sim N(0, \mathcal{I}_1(\theta_0))$

$$\implies \sqrt{n}(\hat{\theta}_n - \theta) \xrightarrow{w} N\left(0, \frac{1}{\mathcal{I}_1(\theta_0)}\right)$$

□

5 Hypothesis Testing

$\underline{X} = (X_1 \dots X_n)$ a random sample from $X \sim f_X(x; \theta)$ $\theta \in \Theta$. We want to use the observed sample to evaluate θ .

First we need a bit of terminology.

Definition 5.0.1. a **statistical test** is defined as a decision rule on the sample space. Using this decision rule and the observed sample we decide to accept or reject an hypothesis about θ .

The hypothesis we decide to accept or refuse is called **the null hypothesis** H_0

Definition 5.0.2. We define a **statistical hypothesis** any sentence that specifies a partition of the statistical model $\{f_x(X; \theta) : \theta \in \Theta\}$.

If the specification of the statistical model is complete then the statistical hypothesis is called **simple** $\theta = \theta^*$.

If the specification is partial is called composite, in particular $\theta > \theta^*, \theta \in R^* \subset \Theta$ is called **composite unidirectional hypothesis**.

We denote by $\omega_0 \subset \Theta$ the set of values that are specified by the null statistical hypothesis H_0 .

We want to check if data support hypothesis of this type:

$$H_0 : \theta \in \Theta$$

$$H_1 : \theta \notin \Theta$$

The null statistical hypothesis is something that exist before the experiment. It is an hypothesis that exist until proved otherwise. The alternative of H_0 (denoted by H_1) is the complement of H_0 .

The decision rule about accepting or rejecting H_0 has different interpretations, in particular:

- if we decide, based on data, to reject H_0 , this has a unique consequence
- if we decide, based on data, to accept H_0 this does not imply the support of H_0 .

The hypothesis test is a decision on the sample space and this decision rule is characterized by a function defined on the set of all the possible value of $(X_1 \dots X_n)$. There will be some points $\underline{X} \in R_0 \subset \mathbb{R}^n$ such that the decision rule leads to reject H_0 and some other points such that the decision rule leads to do not reject H_0 .

Definition 5.0.3. the subspace R_0 is called **rejection region for H_0** .

The decision rule is determined by the following criteria:

- if $x \in R_0$ then reject H_0
- if $x \notin R_0$ do not reject H_0 .

With this space, the decision rule provides a partition of the sample space into two subsets. However we can work with a sub partition of the sample space (\mathbb{R}^n) only if n is small.

To reduce the dimensionality of the decision rule (or the partition of the sample space) we can use a sufficient statistic $T_n = T(\underline{x})$ with $T : \mathbb{R}^n \rightarrow \mathbb{R}^m$ with $m \leq n$. The original decision rule is transformed in the new decision rule:

- If $T_n \in C_0$ then we reject H_0
- if $T_n \notin C_0$ then we do not reject H_0

Where C_0 is our new rejection region.

The statistical hypothesis defines a partition of the parametric space Θ .

The decision rule defines a partition of the sample space (or a lower dimensional counterpart).

So we can see a test as a link between the two partitions.

Confronting the choices done with the test with the real situation it can happen 4 different scenarios:

| | $X \in R_0$ (reject H_0) | $X \notin R_0$ (accept H_0) |
|-------------------------------|-----------------------------|--------------------------------|
| $H_0: \theta \in \omega_0$ | E_1 | G_1 |
| $H_1: \theta \notin \omega_0$ | G_2 | E_2 |

G_1 accept H_0 when H_0 is true, so we took the right decision

G_2 reject H_0 when H_0 is false, so we took the right decision

E_1 reject H_0 when H_0 is true, so we took the wrong decision

E_2 accept H_0 when H_0 is false, so we took the wrong decision

Definition 5.0.4. we define the function $\gamma(\theta)$ with $\theta \in \Theta$ in this way

- if $\theta \in \Theta_0$ $\gamma(\theta) := \mathbb{P}(\underline{X} \in R_0; \theta \in \Theta_0)$
- if $\theta \in \Theta_0^c$ $\gamma(\theta) := \mathbb{P}(\underline{x} \in R_0, \theta \notin \Theta_0)$

Definition 5.0.5. Some terminology.

E_1 is called **error of the first type** and

$$\alpha := \mathbb{P}(E_1) = \mathbb{P}(\text{reject } H_0, H_0 \text{ true}) = \mathbb{P}(\underline{x} \in R_0, \theta \in \omega_0)$$

α is called **size of the test**.

E_2 is called **error of the second type** and

$$\beta := \mathbb{P}(E_2) = \mathbb{P}(\text{accept } H_0, H_0 \text{ false}) = \mathbb{P}(\underline{x} \notin R_0, \theta \notin \omega_0)$$

note that $\gamma = 1 - \beta = \mathbb{P}(\underline{x} \in R_0, \theta \notin \omega_0)$.

$1 - \beta$ is called **power of the test**.

Definition 5.0.6. for $\alpha \in [0, 1]$ a test with power function $\gamma(\theta)$ is a **size α test** is :

$$\sup_{\theta \in \Theta_0} \gamma(\theta) = \alpha$$

Definition 5.0.7. for $\alpha \in [0, 1]$ a test with power function $\gamma(\theta)$ is a **level α test** if

$$\sup_{\theta \in \Theta_0} \gamma(\theta) \leq \alpha$$

We want to find a rejection region R_0 where α and β are small.

But it is impossible to reduce α and at the same time β as shown in the next observation 19

Observation 19. Fix a rejection region C_0 and define

- $\alpha = \alpha(C_0)$
- $\beta = \beta(C_0)$

Suppose that C'_0 is a different reject region such that

$$\alpha(C'_0) < \alpha(C_0)$$

Then we want to show that $\beta(C'_0) > \beta(C_0)$.
Indeed if $\alpha(C'_0) < \alpha(C_0)$ then $\mathbb{P}(C'_0) < \mathbb{P}(C_0)$ because

$$\alpha(C'_0) = \mathbb{P}(\underline{x} \in C'_0 | H_0) < \mathbb{P}(\underline{x} \in C_0 | H_0) = \alpha(C_0)$$

so $\mathbb{P}((C'_0)^c) > \mathbb{P}(C_0^c)$

$$\beta(C'_0) = \mathbb{P}(\underline{x} \notin C'_0 | H_1) \mathbb{P}(\underline{x} \in (C'_0)^c | H_1) > \mathbb{P}(\underline{x} \in C_0^c | H_1) = \mathbb{P}(\underline{x} \notin C_0 | H_1) = \beta(C_0)$$

So we have that if α decrease β increase.

The most dangerous error is the type one error so usually we fix α to control it and look for the reject region C_0 which has α as probability of committing an error of the first type and the smallest possible β as probability of committing an error of the second type.

Example 5.0.1. Consider $\underline{X} = (X_1 \dots X_n)$ a random sample from $X \sim \text{Ber}(\theta)$ (Flip of a coin).
 $\hat{\theta}_n = \frac{1}{n} \sum_{i=1}^n X_i$.

We consider the following test:

$$H_0 : \theta = \frac{1}{2}$$

$$H_1 : \theta = \frac{2}{3}$$

The easiest way to test this hypothesis is by taking an experiment.
Suppose we film the coin 30 times, we have: $\sum_{i=1}^n X_i \sim \text{Bin}(30, \theta)$.
Now guess some reject regions

C'_0 we reject H_0 if the number of heads is strictly greater than 25

C''_0 we reject H_0 if the number of heads is strictly greater than 17

In C'_0 we expect a large β because we will "reject very often".

In C''_0 we expect a large α because we will "reject quite rarely".

- $\alpha(C'_0) = \mathbb{P}(\sum_{i=1}^{30} x_i > 25 | \theta = \frac{1}{2}) = \sum_{M=25}^{30} \binom{30}{M} \left(\frac{1}{2}\right)^M \left(\frac{1}{2}\right)^{30-M} = 0.000029\dots$
- $\beta(C'_0) = \mathbb{P}(\sum_{i=1}^{30} x_i < 25 | \theta = \frac{2}{3}) = \sum_{M=25}^{30} \binom{30}{M} \left(\frac{2}{3}\right)^M \left(\frac{1}{3}\right)^{30-M} = 0.9877\dots$
- $\alpha(C''_0) = \mathbb{P}(\sum_{i=1}^{30} x_i > 17 | \theta = \frac{1}{2}) = \sum_{M=17}^{30} \binom{30}{M} \left(\frac{1}{2}\right)^M \left(\frac{1}{2}\right)^{30-M} = 0.18\dots$
- $\beta(C''_0) = \mathbb{P}(\sum_{i=1}^{30} x_i < 17 | \theta = \frac{2}{3}) = \sum_{M=17}^{30} \binom{30}{M} \left(\frac{2}{3}\right)^M \left(\frac{1}{3}\right)^{30-M} = 0.166\dots$

Definition 5.0.8. Let C be the class of tests for testing $H_0: \theta \in \Theta_0$ versus $H_1: \theta \in \Theta_0^c$.

A test in the class C , with power function $\gamma(\theta)$ is a **uniformly most powerful (UMP)** class C test if

$$\gamma(\theta) \geq \gamma'(\theta)$$

for every $\theta \in \Theta_0^c$ and every γ' that is the power function of a test in the class C .

For us the class C will be the class of tests of size α . Then we will write **UMP level α test**

In most cases this test does not exist. A situation in which this test exist is when H_0, H_1 are simple. In other words the requirements of this definition are very strong, in particular the fact that we are considering arbitrary hypothesis. In order to identify the UMP level α test we restrict to the class of simple hypothesis. So we can use the "Neyman-Pearson lemma".

Theorem 5.0.1. Neyman-Pearson lemma

Consider the problem of testing $H_0 : \theta = \theta_0$ versus $H_1 : \theta = \theta_1$. Let $f_{\underline{X}}(\underline{x}; \theta)$ be the joint density function of $\underline{X} = (X_1 \dots X_n)$. Consider the test with the rejection region R_0 that satisfies

$A: \underline{x} \in R_0$ if $f_{\underline{x}}(\underline{x}; \theta_1) > k f_{\underline{x}}(\underline{x}; \theta_0)$ and $\underline{x} \in R_0^c$ if $f_{\underline{x}}(\underline{x}; \theta_1) < k f_{\underline{x}}(\underline{x}; \theta_0)$ for some $k \geq 0$

$B: \alpha = \mathbb{P}(\underline{x} \in R_0; \theta_0)$

A is called the **shape of the rejection region**

B is called **size of the rejection region**.

Then

- any test that satisfies A and B is a UMP level α test
- if exist a test that satisfies A and B with $k \geq 0$ then every UMP level α test is a size α test (B) and every UMP level α test satisfies A except on a set such that:

$$\mathbb{P}(\underline{x} \in A; \theta_0) = \mathbb{P}(\underline{x} \in A; \theta_1) = 0$$

Example 5.0.2. Let $\underline{X} = (X_1 \dots X_n)$ a random sample from $X \sim N(\theta, \sigma^2)$, where σ^2 is known. We test:

$$H_0 : \theta = \theta_0 \quad vs \quad H_1 : \theta = \theta_1$$

with $\theta_1 > \theta_0$

$$\frac{\mathcal{L}(\theta_1; \underline{x})}{\mathcal{L}(\theta_0; \underline{x})} \geq c$$

with $\alpha = \mathbb{P}(\underline{x} \in R_0; \theta_0)$ fixed

$$\begin{aligned} \implies & \exp \left\{ -\frac{1}{2\sigma^2} \left[\sum_{i=1}^n (x_i - \theta_1)^2 - \sum_{i=1}^n (x_i - \theta_0)^2 \right] \right\} \geq c \\ \iff & \exp \left\{ -\frac{1}{2\sigma^2} \left[\sum_{i=1}^n x_i^2 - n\theta_1^2 - 2\theta_1 \sum_{i=1}^n x_i - \sum_{i=1}^n x_i^2 - n\theta_0^2 + 2\theta_0 \sum_{i=1}^n x_i \right] \right\} \geq c \\ \iff & \exp \left\{ -\frac{1}{2\sigma^2} \left[-n\theta_1^2 - 2\theta_1 \sum_{i=1}^n x_i - n\theta_0^2 + 2\theta_0 \sum_{i=1}^n x_i \right] \right\} \geq c \\ \iff & \exp \left\{ -\frac{1}{2\sigma^2} \left[n(\theta_1^2 - \theta_0^2) - 2 \sum_{i=1}^n x_i(\theta_1 - \theta_0) \right] \right\} \geq c \\ \iff & \exp \left\{ -\frac{1}{2\sigma^2} (\theta_1 - \theta_0) \left[n(\theta_1 + \theta_0) - 2 \sum_{i=1}^n x_i \right] \right\} \geq c \\ \iff & -(\theta_1 - \theta_0) \left[n(\theta_1 + \theta_0) - 2 \sum_{i=1}^n x_i \right] \geq 2 \ln(c) \sigma^2 \\ \iff & n(\theta_1 + \theta_0) - 2 \sum_{i=1}^n x_i \leq -\frac{2 \ln(c) \sigma^2}{\theta_1 - \theta_0} \\ \iff & -2 \sum_{i=1}^n x_i \leq -\frac{2 \ln(c) \sigma^2}{\theta_1 - \theta_0} - n(\theta_1 + \theta_0) \\ \iff & \sum_{i=1}^n x_i \geq -\frac{\ln(c) \sigma^2}{\theta_1 - \theta_0} - \frac{n}{2} (\theta_1 + \theta_0) \end{aligned}$$

So $\bar{X}_n \geq c$ is the shape of the rejection region. Where now c is: $-\frac{\ln(c) \sigma^2}{\theta_1 - \theta_0} - \frac{n}{2} (\theta_1 + \theta_0)$ So to find the

value of c given α we remember that $\bar{X}_n \sim N(\theta, \frac{\sigma}{\sqrt{n}})$ so:

$$\begin{aligned}\alpha &= \mathbb{P}(\bar{X}_n \geq c | H_0 : \theta = \theta_0) \\ &= \mathbb{P}\left(\frac{\bar{X}_n - \theta_0}{\sigma/\sqrt{n}} \geq \frac{c - \theta_0}{\sigma/\sqrt{n}}\right) \\ &= \mathbb{P}\left(Z \geq \frac{c - \theta_0}{\sigma/\sqrt{n}}\right) \\ &= 1 - \Phi\left(\frac{c - \theta_0}{\sigma/\sqrt{n}}\right)\end{aligned}$$

Where Z is the standard Gaussian. In the end $c = \theta_0 + z_\alpha \frac{\sigma}{\sqrt{n}}$.
According to the Neyman-Pearson lemma 5.0.1

$$\left\{\bar{X}_n \geq \theta_0 + z_\alpha \frac{\sigma}{\sqrt{n}}\right\} \quad \text{is the UMP level } \alpha$$

Example 5.0.3. Let $\underline{X} = (X_1 \dots X_n)$ a random sample from $X \sim \text{Ber}(\theta)$. We test:

$$H_0 : \theta = \theta_0 \quad \text{vs} \quad H_1 : \theta = \theta_1$$

With $\theta_1 > \theta_0$

$$\begin{aligned}\Rightarrow \frac{\mathcal{L}(\theta_0; \underline{x})}{\mathcal{L}(\theta_1; \underline{x})} &\leq k \\ \frac{\theta_0^{\sum_{i=1}^n x_i} (1 - \theta_0)^{n - \sum_{i=1}^n x_i}}{\theta_1^{\sum_{i=1}^n x_i} (1 - \theta_1)^{n - \sum_{i=1}^n x_i}} &= \left(\frac{\theta_0}{\theta_1}\right)^{\sum_{i=1}^n x_i} \left(\frac{1 - \theta_0}{1 - \theta_1}\right)^{n - \sum_{i=1}^n x_i} \leq k \\ \ln \left(\frac{\theta_0}{\theta_1}\right) \sum_{i=1}^n x_i + \left(n - \sum_{i=1}^n x_i\right) \ln \left(\frac{1 - \theta_0}{1 - \theta_1}\right) &\leq \ln k \\ \ln \left(\frac{\theta_0}{\theta_1}\right) \sum_{i=1}^n x_i + n \ln \left(\frac{1 - \theta_0}{1 - \theta_1}\right) - \ln \left(\frac{1 - \theta_0}{1 - \theta_1}\right) \sum_{i=1}^n x_i &\leq \ln k \\ \sum_{i=1}^n x_i \left[\ln \left(\frac{\theta_0}{\theta_1}\right) - \ln \left(\frac{1 - \theta_0}{1 - \theta_1}\right) \right] &\leq \ln k - n \ln \left(\frac{1 - \theta_0}{1 - \theta_1}\right)\end{aligned}$$

because $\ln \left(\frac{\theta_0}{\theta_1}\right)$ and $-\ln \left(\frac{1 - \theta_0}{1 - \theta_1}\right)$ are ≤ 0 we have

$$\sum_{i=1}^n x_i \geq k$$

where $\sum_{i=1}^n x_i \sim \text{Bin}(\theta, n)$.

$$\alpha = \mathbb{P}\left(\sum_{i=1}^n x_i \geq k; \theta = \theta_0\right) = \sum_{r=k}^n \binom{n}{r} \theta_0^r (1 - \theta_0)^{n-r}.$$

We have $k \in \mathbb{N}$ so we can not find a k such that $\mathbb{P}\left(\sum_{i=1}^n x_i \geq k; \theta = \theta_0\right)$ is exactly equal to α . So we look for a value that provides a value very close to α

Example 5.0.4. Let $\underline{X} = (X_1 \dots X_n)$ a random sample from $X \sim \text{EF}()$ (Exponential Family) so $f_X(x; \theta) = \exp\{A(x)Q(\theta) + C(x) - K(\theta)\}$. Assume Q monotone increasing. We want to test

$$H_0 : \theta = \theta_0 \quad \text{vs} \quad H_1 : \theta = \theta_1$$

With $\theta_0 < \theta_1$. Apply the Neyman-Pearson lemma 5.0.1

$$\begin{aligned}
&\implies \frac{\mathcal{L}(\theta_0; \underline{x})}{\mathcal{L}(\theta_1; \underline{x})} \leq k \\
&\iff \frac{\prod_{i=1}^n \exp\{A(x_i)Q(\theta_0) + C(x_i) - K(\theta_0)\}}{\prod_{i=1}^n \exp\{A(x_i)Q(\theta_1) + C(x_i) - K(\theta_1)\}} \leq k \\
&\iff \exp\{Q(\theta_0) \sum_{i=1}^n A(x_i) + \sum_{i=1}^n C(x_i) - nK(\theta_0) - Q(\theta_1) \sum_{i=1}^n A(x_i) + \sum_{i=1}^n C(x_i) + nK(\theta_1)\} \leq k \\
&\iff \exp\{Q(\theta_0) \sum_{i=1}^n A(x_i) - nK(\theta_0) - Q(\theta_1) \sum_{i=1}^n A(x_i) + nK(\theta_1)\} \leq k \\
&\iff \exp\{\sum_{i=1}^n A(x_i)(Q(\theta_0) - Q(\theta_1)) - n(K(\theta_0) - K(\theta_1))\} \leq k \\
&\iff \sum_{i=1}^n A(x_i)(Q(\theta_0) - Q(\theta_1)) - n(K(\theta_0) - K(\theta_1)) \leq \ln k \\
&\iff \sum_{i=1}^n A(x_i)(Q(\theta_0) - Q(\theta_1)) \leq c \\
&\iff \sum_{i=1}^n A(X_i) \geq c
\end{aligned}$$

where in the last passage we used $(Q(\theta_0) - Q(\theta_1)) < 0$ because $\theta_1 > \theta_0$ and Q is strictly increasing. $\sum_{i=1}^n A(X_i) \geq c$ is the shape of the reject region for an arbitrary exponential family with Q strictly increasing

Proof. of the Neyman-Pearson lemma 5.0.1 Any test that satisfies B is a size α test, and hence a level α test because

$$\sup_{\theta \in \Theta_0} \mathbb{P}(\underline{x} \in R) = \alpha$$

Consider a test function in the sample space that is 1 if $\underline{x} \in R$ and 0 if $\underline{x} \in R^C$.

- let $\Phi(\underline{x})$ be a test function that satisfies A and B
- let $\Phi'(\underline{x})$ be a test function (dichotomic) of any other level α test.
- let $\gamma(\theta)$ and $\gamma'(\theta)$ be the corresponding power functions (when $\theta \in \Theta_0^c$)

Observe that $\Phi(\underline{x}) = 1$ if $\underline{x} \in R$ if $f_{\underline{X}}(\underline{x}; \theta_1) > k f_{\underline{X}}(\underline{x}; \theta_0)$

$\Phi(\underline{x}) = 0$ if $\underline{x} \in R^c$ if $f_{\underline{X}}(\underline{x}; \theta_1) < k f_{\underline{X}}(\underline{x}; \theta_0)$

From this we get

$$(\Phi(\underline{x}) - \Phi'(\underline{x}))(f_{\underline{X}}(\underline{x}; \theta_1) - k f_{\underline{X}}(\underline{x}; \theta_0)) \geq 0$$

Then

$$\int (\Phi(\underline{x}) - \Phi'(\underline{x}))(f_{\underline{X}}(\underline{x}; \theta_1) - k f_{\underline{X}}(\underline{x}; \theta_0)) = \gamma(\theta_1) - \gamma'(\theta_1) - k[\gamma(\theta_0) - \gamma(\theta_1)] \geq 0$$

Now observe that Φ' is a level α test and ϕ is a size α test, this implies that

$$0 \leq \gamma(\theta_1) - \gamma'(\theta_1) - k\gamma(\theta_0) - \gamma'(\theta_0) \leq \gamma(\theta_1) - \gamma'(\theta_1)$$

$$\implies 0 \leq \gamma(\theta_1) - \gamma'(\theta_1)$$

$\implies \gamma(\theta_1) \geq \gamma'(\theta_1)$ We proved the test defined by Neyman-Pearson condition A, B is a UMP level α test.

Second part

Let Φ' be a test function for every UMP level α test. By the previous part, that is the test Φ satisfying

$A; B$ is also a UMP level α test. This implies

$$\gamma(\theta_1) = \gamma'(\theta_1)$$

This implies that for any $k > 0$

$$\alpha - \gamma'(\theta_0) = \gamma(\theta_0) - \gamma'(\theta_1) \leq 0$$

Now since Φ' is a level α test we have that

$$\gamma'(\theta_0) \leq \alpha \implies \gamma'(\theta_0) = \alpha$$

This implies that Φ' is a size α test. Finally Recall that

$$(\Phi(\underline{x}) - \Phi'(\underline{x}))(f_{\underline{X}}(\underline{x}; \theta_1) - kf_{\underline{X}}(\underline{x}; \theta_0)) \geq 0$$

if we integrate over the sample space, we have 0 only if Φ' satisfies assumption A. Except on a set A of probability 0:

$$\int_A f_{\underline{X}}(\underline{x}; \theta_1) d\mathbf{x} = 0 \quad i = 0, 1$$

□

We will see extensions of the Neyman-Pearson Lemma for

- unidirectional hypothesis (robin test)
- arbitrary hypothesis (likelihood ratio test)

Corollary 5.0.1. *(link between Neyman-Pearson Lemma and sufficiency)*

Consider the hypothesis of the Neyman-Pearson Lemma 5.0.1. Suppose $T(\underline{X})$ a sufficient statistic for the parameter θ . Denote by $f_{\underline{X}}(\underline{x}; \theta_i)$ for $i = 0, 1$ the joint density distribution function of the random sample and by $g(t; \theta_i)$, $i = 0, 1$ the density function of the sufficient statistic.

Then any test based on T whit rejection region S (a subset o the sample space defined by T) is a UMP level α test if it satisfies the two following:

- $t \in S$ if $g(t; \theta_1) > kg(t; \theta_0)$
- $t \in S^c$ if $g(t; \theta_1) < kg(t; \theta_0)$

for some $k \geq 0$, where $\alpha = \mathbb{P}(T \in S; \theta_0)$

Proof. In terms of the random sample $(X_1 \dots X_n)$, the test based on the sufficient static T has the following rejection region

$$R = \{\underline{x} : T(\underline{x}) \in S\}$$

By the factorization theorem the distribution of the random sample $(X_1 \dots X_n)$ can be written as

$$f_{\underline{X}}(\underline{x}; \theta_i) = g(T(\underline{x}); \theta_i)h(\underline{x}) \quad i = 0, 1$$

Consider the condition $t \in S$ if $g(t; \theta_1) > g(t; \theta_0)$ and multiply both sides of the inequality by the non negative function h . We see that the rejection region R satisfies $\underline{x} \in R$ if

$$f_{\underline{X}}(\underline{x}; \theta_1) = g(T(\underline{x}); \theta_1)h(\underline{x}) < g(T(\underline{x}); \theta_0)h(\underline{x}) = f_{\underline{X}}(\underline{x}; \theta_0)$$

By the size condition we have

$$\mathbb{P}(\underline{X} \in R; \theta_0) = \mathbb{P}(T(\underline{X}) \in S; \theta_0) = \alpha$$

At this point we remember the Neyman-Pearson Lemma 5.0.1 and the test based on the sufficient statistic T is a UMP level α test. □

Example 5.0.5. Remember the example 5.0.4 of with the exponential family $f_X(x; \theta) = \exp\{A(x)Q(\theta) + C(x) - K(\theta)\}$ we know a sufficient statistic is $\sum_{i=1}^n A(X_i)$ We want to test

$$H_0 : \theta = \theta_0 \quad vs \quad H_1 : \theta = \theta_1$$

Assume Q monotone increasing and $\theta_1 > \theta_0$. So $Q(\theta_1) > Q(\theta_0)$.

The shape of the rejection region is

$$\sum_{i=1}^n A(x_i) \geq c$$

which is a function of a sufficient statistic for θ

In general, by the previous corollary 5.0.1 the shape of the reject region defined by the Neyman-Pearson Lemma is a function of the sufficient statistic for the parameter of interest.

6 Exercises

Exercise 1. Let X_1 and X_2 two random variables independent and uniformly distributed on the interval $[0, 1]$. Find the distribution of:

1. $Y = X_1 + X_2$
2. $W = \frac{X_1}{X_2}$
3. $Z = X_1 X_2$

Solution:

1. Let's start with the sum of two generic random variables:
we know that $f_{Y|X_1}(y|x_1) = f_{X_2}(y - x_1)$ and the joint distribution on two random variables is:

$$f_{X_1, Y}(x_1, y) = f_{Y|X_1}(y|x_1)f_{X_1}(x_1)$$

so in our case:

$$f_{X_1, Y}(x_1, y) = f_{X_2}(y - x_1)f_{X_1}(x_1)$$

and now we can calculate the PDF of Y simply by integrating the PDF of the joint distribution:

$$\begin{aligned} f_Y(y) &= \int_{-\infty}^{\infty} f_{X_1, Y}(x_1, y) dx_1 \\ &= \int_{-\infty}^{\infty} f_{X_2}(y - x_1)f_{X_1}(x_1) dx_1 \end{aligned}$$

(For a more detailed analysis see the convolution product).

Now we can proceed replacing the generic PDF with the one of a uniform distribution on the interval $[0, 1]$ is $\mathbb{1}_{[0,1]}(t)$, we have:

$$\begin{aligned} f_Y(y) &= \int_{-\infty}^{\infty} \mathbb{1}_{[0,1]}(y - x_1) \mathbb{1}_{[0,1]}(x_1) dx_1 \\ &= \int_0^1 \mathbb{1}_{[0,1]}(y - x_1) dx_1 \end{aligned}$$

and by separating the integral in various cases we can solve it obtaining:

$$f_Y(y) = \begin{cases} 0 & \text{if } y < 0 \\ y & \text{if } y \in [0, 1] \\ 2 - y & \text{if } y \in (1, 2] \\ 0 & \text{if } y > 2 \end{cases}$$

2. for the distribution of W we will use the CDF function:
first of all it is easy to prove that for $w \leq 0$, $F_W(w) = 0$, so in the next passage we can assume

$$w > 0$$

$$\begin{aligned}
F_W(w) &= \mathbb{P}[W < w] \\
&= \mathbb{P}\left[\frac{X_2}{X_1} < w\right] \\
&= \mathbb{P}\left[\frac{X_2}{X_1} < w, X_1 > 0\right] + \mathbb{P}\left[\frac{X_2}{X_1} < w, X_1 < 0\right] \\
&= \mathbb{P}[X_2 < X_1 w, X_1 > 0] + 0 \\
&= \int_0^\infty f_{x_1}(x_1) \int_{-\infty}^{wx_1} f_{x_2}(x_2) dx_2 dx_1 \\
&= \int_0^\infty \mathbb{1}_{[0,1]}(x_1) \int_{-\infty}^{wx_1} \mathbb{1}_{[0,1]}(x_2) dx_2 dx_1 \\
&= \int_0^1 \begin{cases} 1 & \text{if } wx_1 > 1 \\ wx_1 & \text{if } wx_1 < 1 \end{cases} dx_1 \\
&= \begin{cases} \int_0^1 wx_1 dx & \text{if } w \leq 1 \\ \int_0^{\frac{1}{w}} wx_1 dx + \int_{\frac{1}{w}}^1 1 dx_1 & \text{if } w > 1 \end{cases} \\
&= \begin{cases} \frac{w}{2} & \text{if } 0 \leq w \leq 1 \\ 1 - \frac{1}{2w} & \text{if } w > 1 \end{cases}
\end{aligned}$$

3. for the distribution of Z we will adopt a more straight forward approach and we will use the CDF function:

$$\begin{aligned}
F_Z(z) &= \mathbb{P}[Z < z] \\
&= \mathbb{P}[X_1 X_2 < z] \\
&= \mathbb{P}[X_1 X_2 < z, X_1 > 0] + \mathbb{P}[X_1 X_2 < z, X_1 < 0] \\
&= \mathbb{P}[X_2 < \frac{z}{X_1}, X_1 > 0] + 0 \\
&= \int_0^\infty f_{X_1}(x_1) \int_{-\infty}^{z/x_1} f_{X_2}(x_2) dx_2 dx_1 \\
&= \int_0^\infty \mathbb{1}_{[0,1]}(x_1) \int_{-\infty}^{z/x_1} \mathbb{1}_{[0,1]}(x_2) dx_2 dx_1 \\
&= \int_0^\infty \mathbb{1}_{[0,1]}(x_1) \begin{cases} 1 & \text{if } \frac{z}{x_1} > 1 \\ \frac{z}{x_1} & \text{if } \frac{z}{x_1} < 1 \end{cases} dx_1 \\
&= \int_0^1 \begin{cases} 1 & \text{if } \frac{z}{x_1} > 1 \\ \frac{z}{x_1} & \text{if } \frac{z}{x_1} < 1 \end{cases} dx_1 \\
&= \begin{cases} \int_0^z 1 dx_1 + \int_z^1 \frac{z}{x_1} dx_1 & \text{if } z < 1 \\ \int_0^1 1 dx_1 & \text{if } z \geq 1 \end{cases} \\
&= \begin{cases} z - z \ln(z) & \text{if } 0 < z < 1 \\ 1 & \text{if } z \geq 1 \end{cases} .
\end{aligned}$$

Exercise 2. Consider X_1, X_2, X_3 three random variables independent and identically distributed with distribution $\sim \text{Exp}(\frac{1}{2})$.

Find the distribution of:

$$1. U = \frac{X_2}{X_1}$$

$$2. W = \sum_{i=1}^3 X_i$$

(in this exercise we will consider the PDF of the exponential $f_X = \lambda e^{-\lambda x}$).

Solution:

1. For the distribution of U we will use the CDF function:

first of all it is easy to prove that for $u \leq 0$, $F_U(u) = 0$, so in the next passage we can assume $u > 0$

$$\begin{aligned}
 F(u) &= \mathbb{P}[U < u] \\
 &= \mathbb{P}\left[\frac{X_2}{X_1} < u\right] \\
 &= \mathbb{P}\left[\frac{X_2}{X_1} < u, X_1 > 0\right] + \mathbb{P}\left[\frac{X_2}{X_1} < u, X_1 < 0\right] \\
 &= \mathbb{P}[X_2 < X_1 u, X_1 > 0] + 0 \\
 &= \int_0^\infty f_{x_1}(x_1) \int_{-\infty}^{ux_1} f_{x_2}(x_2) dx_2 dx_1 \\
 &= \int_0^\infty \lambda e^{-\lambda x_1} \int_0^{ux_1} \lambda e^{-\lambda x_2} dx_2 dx_1 \\
 &= \int_0^\infty \lambda e^{-\lambda x_1} (1 - e^{-\lambda u x_1}) x_1 \\
 &= 1 - \frac{1}{1+u}
 \end{aligned}$$

2. Here we will use the moment-generating function to demonstrate that $W \sim \text{Gamma}(\alpha = 3, \beta = \frac{1}{2})$.

The the moment-generating function of the exponential with parameter 2 is: $M_X(t) = \frac{1}{\frac{1}{2}-t}$.

$$\begin{aligned}
 M_W(t) &= \mathbb{E}[e^{wt}] \\
 &= \mathbb{E}[e^{\sum_{i=1}^3 X_i t}] \\
 &= \mathbb{E}\left[\prod_{i=1}^3 e^{X_i t}\right] \\
 &= \prod_{i=1}^3 \mathbb{E}[e^{X_i t}] \\
 &= \prod_{i=1}^3 \frac{1}{\frac{1}{2}-t} \\
 &= \left(\frac{1}{\frac{1}{2}-t}\right)^3 \\
 &= \left(\frac{\frac{1}{2}-t}{\frac{1}{2}}\right)^{-3} \\
 &= \left(1 - \frac{t}{\frac{1}{2}}\right)^{-3}
 \end{aligned}$$

which is the MGF of a $\text{Gamma}(3, \frac{1}{2})$

Exercise 3. Let (X, Y) be a bivariate random variablesuch that $X \sim U(-1, 1)$ and $Y|X \sim U(x, x+1)$. Find he distribution of $Z = -\ln(Y - X)$

Solution:

$$\begin{aligned}
 \mathbb{P}(Z < z) &= \mathbb{P}(-\log(X - Y) < t) \\
 &= \mathbb{P}(\log(X - Y) \geq z) \\
 &= \mathbb{P}(Y \geq X + e^{-z}) \\
 &= \int_{-1}^1 \int_{x+e^{-z}}^{x+1} \frac{1}{2} dx dy \\
 &= 1 - e^{-z}
 \end{aligned}$$

Which is the distribution function of neg exp

Exercise 4. Let A, B be two independent and identically distributed random variables with distribution $\sim U(0, h)$. Compute the probability that the equation $Z^2 - 2AZ + B = 0$ doesn't admit real solutions.

Solution: We are asked to compute $\mathbb{P}(A^2 - B < 0)$.

Exercise 5. Let $X \sim \text{Gamma}(r, 1), Y \sim \text{Gamma}(s, 1)$ independent random variables. Find the distribution of

1. $W := X + Y$
2. $Z := \frac{X}{W}$
3. (Z, w)

Solution: Did in class.

Exercise 6. Let X_1, X_2, X_3 random variables with distribution:

- $X_1 \sim \text{Gamma}(\alpha_1, 1)$
- $X_2 \sim \text{Gamma}(\alpha_2, 1)$
- $X_3 \sim \text{Gamma}(\alpha_3, 1)$

Define:

- $Z = \frac{X_1}{X_1 + X_2 + X_3}$
- $W = \frac{X_2}{X_1 + X_2 + X_3}$

Find the distribution of (Z, W)

For this exercise we will use the theorem on page 165 of [1]

Solution: To have a $(3, 3)$ parametrization we will add another random variable $S := X_1 + X_2 + X_3$.

Consider the parametrization

- $X_1 = ZS$
- $X_2 = WS$
- $X_3 = S - S(Z + W)$

The Jacobian matrix is defined as:

$$J := \begin{bmatrix} \frac{\partial \phi_1}{\partial x_1} & \cdots & \frac{\partial \phi_1}{\partial x_n} \\ \vdots & \ddots & \vdots \\ \frac{\partial \phi_N}{\partial x_1} & \cdots & \frac{\partial \phi_N}{\partial x_n} \end{bmatrix}$$

in our situation then:

$$|J| = \left| \begin{bmatrix} S & 0 & Z \\ 0 & S & W \\ -S & -S & 1 - Z - W \end{bmatrix} \right| = S^2$$

So by the previous theorem we have:

$$f_{Z,W,S}(z, w, s) = f_{X_1, X_2, X_3}(zs, ws, s - s(z + w))|J|$$

Remembering that X_1, X_2, X_3 are independent then $f_{X_1, X_2, X_3}(zs, ws, s - s(z + w)) = f_{X_1}(zs)f_{X_2}(ws)f_{X_3}(s - s(z + w))$ so

$$\begin{aligned} f_{Z,W,S}(z, w, s) &= f_{X_1}(zs)f_{X_2}(ws)f_{X_3}(s - s(z + w))s^2 \\ &= \frac{(zs)^{\alpha_1-1}(ws)^{\alpha_2-1}(s - s(z + w))^{\alpha_3-1}e^{-s}}{\Gamma(\alpha_1)\Gamma(\alpha_2)\Gamma(\alpha_3)}s^2\mathbb{1}_{[\min(zs, ws, s - s(z + w)), \infty)}(0) \\ &= \frac{z^{\alpha_1-1}w^{\alpha_2-1}(1 - (z + w))^{\alpha_3-1}}{\Gamma(\alpha_1)\Gamma(\alpha_2)\Gamma(\alpha_3)}s^{\alpha_1+\alpha_2+\alpha_3-1}e^{-s} \end{aligned}$$

Notice that the second member is the kernel of a Gamma distribution with parameters $\alpha_1 + \alpha_2 + \alpha_3, 1$. We know that $s \geq 0$ so $\min(zs, ws, s - s(z + w))$ has the same sign of $\min(z, w, 1 - (z + w))$. To get the distribution of Z, W we have to integrate $f_{Z,W,S}(z, w, s)$ with respect to s :

$$\begin{aligned} f_{Z,W}(z, w) &= \int_0^\infty f_{Z,W,S}(z, w, s)\mathbb{1}_{[\min(z, w, 1 - (z + w)), \infty)}(0)ds \\ &= \mathbb{1}_{[\min(z, w, 1 - (z + w)), \infty)}(0) \int_0^\infty \frac{z^{\alpha_1-1}w^{\alpha_2-1}(1 - (z + w))^{\alpha_3-1}}{\Gamma(\alpha_1)\Gamma(\alpha_2)\Gamma(\alpha_3)}s^{\alpha_1+\alpha_2+\alpha_3-1}e^{-s}ds \\ &= \mathbb{1}_{[\min(z, w, 1 - (z + w)), \infty)}(0) \frac{z^{\alpha_1-1}w^{\alpha_2-1}(1 - (z + w))^{\alpha_3-1}}{\Gamma(\alpha_1)\Gamma(\alpha_2)\Gamma(\alpha_3)} \int_0^\infty s^{\alpha_1+\alpha_2+\alpha_3-1}e^{-s}ds \\ &= \frac{z^{\alpha_1-1}w^{\alpha_2-1}(1 - (z + w))^{\alpha_3-1}}{\Gamma(\alpha_1)\Gamma(\alpha_2)\Gamma(\alpha_3)}\Gamma(\alpha_1 + \alpha_2 + \alpha_3)\mathbb{1}_{[\min(z, w, 1 - (z + w)), \infty)}(0) \end{aligned}$$

Exercise 7. Show that the moment generating function of a random variable $X \sim NEF(\nu)$ is

$$\mathbb{E}[e^{sx}] = e^{K(s+\nu) - K\nu}$$

Solution:

$$\begin{aligned} e[e^{sx}] &= \int e^{sx}e^{x\nu+C(x)-K(\nu)}dx = \int e^{x(s+\nu)+C(x)+K(s+\nu)-K(s+\nu)-K(\nu)}dx \\ &= e^{K(s+\nu)-K(\nu)} \int e^{x(s+\nu)+C(x)-K(s+\nu)}dx \\ &= e^{K(s+\nu)-K(\nu)}1 \\ &= e^{K(s+\nu)-K(\nu)} \end{aligned}$$

Exercise 8. We want to estimate the proportion θ of individuals in a population for which a certain feature X takes value in A . We take a sample from the population of size n and we measure the feature X . Let Z_n be the number of individuals with feature X in A .

- Is $\frac{Z_n}{n}$ unbiased for θ ? Is it consistent? Is asymptotically Gaussian?
- Find a maximum likelihood estimator for θ .

Exercise 9. Consider a finite population of individuals such that only 40% of individuals survive after one week. After one week we check the population and we find r individuals. We want to estimate the size of population.

Exercise 10. (X_1, X_2, \dots, X_n) from X ($f_x(x, \theta)$ where $f_x(x, \theta) = 2\theta x \exp\{-\theta x^2\}1_{\mathbb{R}}(x)$, $\theta > 0$).

- Find the ML estimator for θ
- Find a sufficient statistic for θ

- *Find the moment estimator for θ*
- *Compare the estimators.*

Ringraziamenti

Special thanks to Moritz And Rade.



(Sorry if you printed this page)

Bibliography

- [1] Casella and Berger. *Statistical Inference*. Duxbury press, second edition edition, 2008.